

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and data manipulation
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

0.1 S&DS 355/555: Assignment 1

0.1.1 NetID: sa857

Due: Sep 17, 2019 11:59pm

0.1.2 Imports and settings

```
In [317]: %reset -f
          %matplotlib inline
          import pandas as pd
          import numpy as np
          import matplotlib.pyplot as plt
          import seaborn as sns
```

```
In [318]: # settings
          pd.set_option('display.max_columns', 999)
```

1 Problem 1: Simple Linear Regression (25 points)

1.1 Problem 1.a:

In class we considered linear regression with the model

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

where $\epsilon \sim N(0, \sigma^2)$ for $i = 1, 2, \dots, n$. Suppose that we believe that the true value of β_0 is zero. In this case we now consider the simpler model $Y_i = \beta_1 X_i + \epsilon_i$. Find an expression for $\hat{\beta}_1$, the estimate of β_1 that minimizes the sum of squared residuals for this simpler model

1.1.1 Answer

Usually, $RSS = (y_1 - \hat{\beta}_0 - \hat{\beta}_1 x_1)^2 + (y_2 - \hat{\beta}_0 - \hat{\beta}_1 x_2)^2 + \dots + (y_n - \hat{\beta}_0 - \hat{\beta}_1 x_n)^2$

Now, $RSS = (y_1 - \hat{\beta}_1 x_1)^2 + (y_2 - \hat{\beta}_1 x_2)^2 + \dots + (y_n - \hat{\beta}_1 x_n)^2$

$$RSS = \sum_{i=1}^n (y_i - \hat{\beta}_1 x_i)^2$$

$$\frac{\partial RSS}{\partial \hat{\beta}_1} = \sum_{i=1}^n 2(y_i - \hat{\beta}_1 x_i)(-x_i) = 0$$

$$\sum_{i=1}^n (y_i - \hat{\beta}_1 x_i)(x_i) = 0$$

$$\sum_{i=1}^n x_i y_i = \sum_{i=1}^n \hat{\beta}_1 (x_i)^2$$

$$\frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n (x_i)^2} = \hat{\beta}_1$$

1.2 Problem 1.b:

Download the `fatherson.csv` file on Canvas with the Jupyter notebook for this homework assignment. This dataset, collected by Galton, contains the height of sons and the height of their father. To read it in, use the function below:

```
x = pd.read_csv("fatherson.csv")
```

After reading in this dataset, create a scatterplot of the sons' heights (on the Y-axis) versus the fathers' heights. Use your answer from (a) to calculate the slope of the least-squares line under the model with no intercept:

$$\text{Son}_i = \beta_1 \text{Father}_i + \epsilon_i$$

Add the fitted line to the scatterplot.

1.2.1 Answer

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMIN
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of column
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments


```
In [319]: # read in data
fsData = pd.read_csv("fatherson.csv")

# scatter plot
plt.scatter(fsData.fheight, fsData.sheight)
plt.title("Scatter plot of sons vs fathers heights")
plt.xlabel("Fathers' heights")
plt.ylabel("Sons' heights")

# calculate beta
beta_1 = sum(fsData.fheight.values * fsData.sheight.values) / sum((fsData.fheight.values)**2)
print(f"Beta 1: {beta_1}")

# plot line
plt.plot(fsData.fheight, beta_1 * fsData.fheight, "r")
plt.show()

Beta 1: 1.0139079627134635
```



1.3 Problem 1.c:

Interpret the meaning of the coefficient β_1 in the context of Galton's father-son dataset.

1.3.1 Answer

Every unit increase in the height of a father leads to a 1.014 increase in the height of his son. In other words, the son's height is a constant multiple of the father's height.

1.4 Problem 1.d:

Use the equations provided in class (for the least-squares coefficients of the linear regression model that includes an intercept) to calculate the least-squares estimates of the coefficients for the linear model that includes a slope and an intercept:

$$\text{Son}_i = \beta_0 + \beta_1 \text{Father}_i + \epsilon_i$$

1.4.1 Answer

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

```
In [320]: # calculate estimates
yBar = np.mean(fsData.sheight)
xBar = np.mean(fsData.fheight)
xSubXBar = fsData.fheight.values - xBar
ySubYBar = fsData.sheight.values - yBar
beta1 = sum(xSubXBar * ySubYBar) / sum(xSubXBar ** 2)
beta0 = yBar - (beta1 * xBar)
print(f"B1: {beta1}, B0: {beta0}")

B1: 0.5140930386233066, B0: 33.886604354077996
```


The estimates are: B1: 0.5140930386233066, B0: 33.886604354077996

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and classification
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [321]: # plot to test
plt.scatter(fsData.fheight, fsData.sheight)
plt.title("Scatter plot of sons vs fathers heights")
plt.xlabel("Fathers' heights")
plt.ylabel("Sons' heights")
plt.plot(fsData.fheight, (beta1*fsData.fheight) + beta0, "r")
```

```
Out[321]: <matplotlib.lines.Line2D at 0x12ee7c510>
```



2 Problem 2: Linear regression and classification (30 points)

Citi Bike is a public bicycle sharing system in New York City. There are hundreds of bike stations scattered throughout the city. Customers can check out a bike at any station and return it at any other station. Citi Bike caters to both commuters and tourists. Details on this program can be found at <https://www.citibikenyc.com/> (<https://www.citibikenyc.com/>)

For this problem, you will build models to predict Citi Bike usage, in number of trips per day. The dataset consists of Citi Bike usage information and weather data recorded from Central Park.

In the `citibike_*.csv` files, we see:

1. `date`
2. `trips`: the total number of Citi Bike trips. This is the outcome variable.
3. `n_stations`: the total number of Citi Bike stations in service
4. `holiday`: whether or not the day is a work holiday
5. `month`: taken from the date variable
6. `dayofweek`: taken from the date variable

In the `weather.csv` file, we have:

1. `date`
2. `PRCP`: amount precipitation (i.e. rainfall amount) in inches
3. `SNWD`: snow depth in inches
4. `SNOW`: snowfall in inches
5. `TMAX`: maximum temperature for the day, in degrees F
6. `TMIN`: minimum temperature for the day, in degrees F
7. `AWND`: average windspeed

You are provided a training set consisting of data from 7/1/2013 to 3/31/2016, and a test set consisting of data after 4/1/2016. The weather file contains weather data for the entire year.

2.1 Problem 2.a: Read in and merge the data.

To read in the data, you can run, for example:

```
train = pd.read_csv("citibike_train.csv")
test = pd.read_csv("citibike_test.csv")
```

Merge the training and test data with the weather data, by date. Once you have successfully merged the data, you may drop the "date" variable; we will not need it for the rest of this assignment.

2.1.1 Import data

```
In [322]: train = pd.read_csv("citibike_train.csv")
test = pd.read_csv("citibike_test.csv")
weather = pd.read_csv("weather.csv")
```

2.1.2 Merge data

```
In [323]: # merge the weather by "date", and then drop "date" column
train = train.merge(weather, left_on="date", right_on="date").drop(columns=["date"])
test = test.merge(weather, left_on="date", right_on="date").drop(columns=["date"])
```

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regressi
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and
 - ▼ 2.1 Problem 2.a: Read in and merge
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regressor
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of column:
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classificati
 - 2.3.1 Compute binary Y, find me
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Ga
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and impo
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator fu
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments


For the rest of this problem, you will train your models on the training data and evaluate them on the test data.

As always, before you start any modeling, you should look at the data. Make scatterplots of some of the numeric variables. Look for outliers and strange values. Comment on any steps you take to remove entries or otherwise process the data. Also comment on whether any predictors are strongly correlated with each other.

2.1.3 Explore data

2.1.3.1 Scatterplot of TMIN vs TMAX

```
In [324]: # make scatterplots
plt.scatter(train.TMIN, train.trips)
plt.scatter(train.TMAX, train.trips)
plt.xlabel("TMIN in blue, TMAX in orange")
plt.ylabel("No. of trips")
plt.show()
```



TMAX and TMIN are strongly correlated with each other. As temperatures rise, the number of train trips increase. TMAX and TMIN are just cutoffs for the same temperature rise phenomenon.

2.1.3.2 Pair plot of all numerical variables

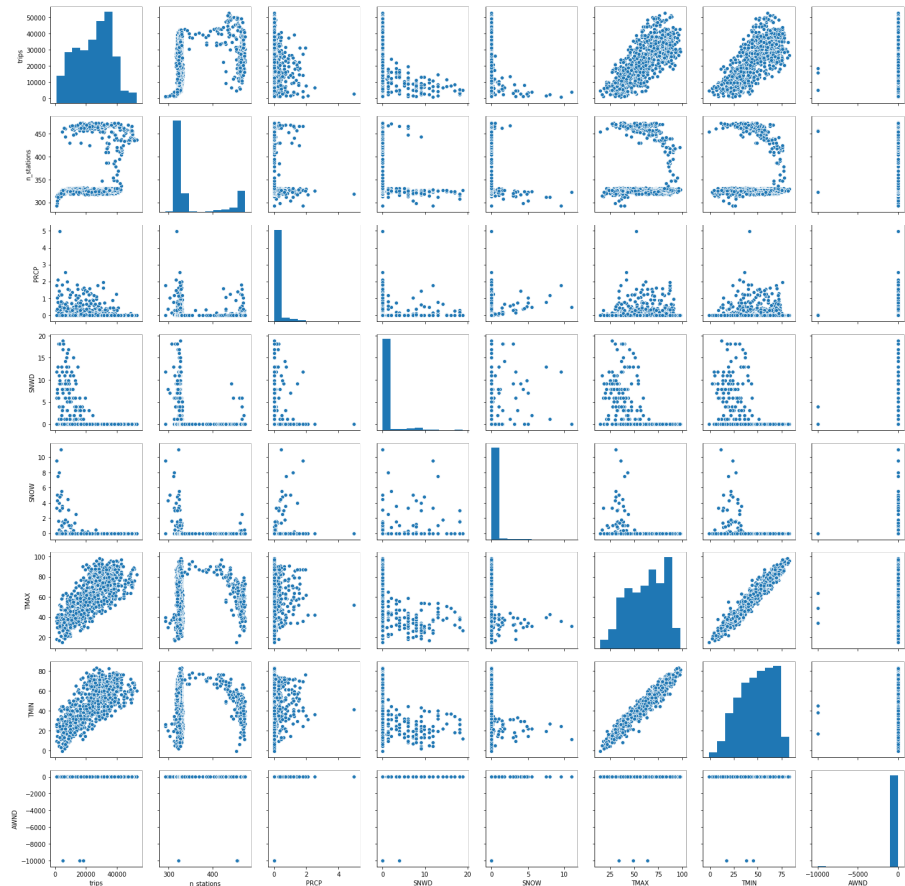
The pair plot below shows strong correlations between "TMAX" and "TMIN", and between "trips" and each of "TMAX" and "TMIN". There is perhaps weak negative correlation between "trips" and "PRCP"

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and preprocessing
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMIN
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [325]: pairPlotDf = train.drop(columns=["holiday", "month", "dayofweek"])
sns.pairplot(pairPlotDf)
```

```
Out[325]: <seaborn.axisgrid.PairGrid at 0x12b5d0a10>
```



2.2 Problem 2.b: Linear regression

Fit a linear regression model to predict the number of trips. Include all the covariates in the data. You may import the `statsmodels.api` module to get a R-like statistical output. You may write code as:

```
import statsmodels.api as sm
X = sm.add_constant(X) # to get the intercept term
model = sm.OLS(y,X).fit()
model.summary()
```

Next, find the "best" linear model that uses only p variables, for each $p = 1, 2, 3, 4, 5$. It is up to you to choose how to select the "best" subset of variables. (A categorical variable or factor such as "month" corresponds to a single variable.) Describe how you selected each model. Give the R^2 and the mean squared error (MSE) on the training and test set for each of the models. Which model gives the best fit to the data? Comment on your findings.

2.2.1 Preprocessing

2.2.1.1 Categorical variables

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regressi
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and
 - ▼ 2.1 Problem 2.a: Read in and merge
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regressor
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of column:
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p == 5$
 - 2.2.3.3 Best model $p == 4$
 - 2.2.3.4 Best model $p == 3$
 - 2.2.3.5 Best model $p == 2$
 - 2.2.3.6 Best model $p == 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find me
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gai
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and impo
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator fu
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [326]: # PROCESSING

# if statement to run the conversion once only
if 'holiday' in train.columns:
    train['holiday'] = train['holiday'].astype('category')
    test['holiday'] = test['holiday'].astype('category')

# put dummies
train = pd.get_dummies(train)
test = pd.get_dummies(test)

# see if test is missing columns
one = set(test.columns.values)
two = set(train.columns.values)
missing_columns = list(two.difference(one))

# fill in columns of zeros to get train and test to be same shape
for col in missing_columns:
    test[col] = 0
```

2.2.1.2 Reordering of columns

```
In [327]: # sort to be the same column order
test = test.reindex(sorted(test.columns), axis=1)
train = train.reindex(sorted(train.columns), axis=1)
```

2.2.2 Model all covariates

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs dayofweek
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gas Station
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

In [328]: # fit a linear regression model using all covariates

```
import statsmodels.api as sm
covariates = train.drop(columns="trips")
X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
allCovariatesResult = model.fit()

print(allCovariatesResult.summary())
print(f"\n\nMSE: {allCovariatesResult.mse_resid}")
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.873
Model:	OLS	Adj. R-squared:	0.870
Method:	Least Squares	F-statistic:	268.5
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	0.00
Time:	22:08:45	Log-Likelihood:	-9747.3
No. Observations:	1001	AIC:	1.955e+04
Df Residuals:	975	BIC:	1.967e+04
Df Model:	25		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.255e+04	754.560	-16.636	0.000	-1.4e+04	-1.11e+04
AWND	0.7647	0.244	3.138	0.002	0.287	1.243
PRCP	-8214.3391	396.200	-20.733	0.000	-8991.842	-7436.836
SNOW	2.3998	188.650	0.013	0.990	-367.808	372.607
SNWD	-215.1752	62.855	-3.423	0.001	-338.521	-91.829
TMAX	352.8827	29.709	11.878	0.000	294.581	411.184
TMIN	-73.3230	32.807	-2.235	0.026	-137.704	-8.942
dayofweek_Fri	-309.4888	337.238	-0.918	0.359	-971.285	352.308
dayofweek_Mon	-1039.1399	341.749	-3.041	0.002	-1709.788	-368.492
dayofweek_Sat	-5414.7253	341.635	-15.849	0.000	-6085.150	-4744.301
dayofweek_Sun	-6381.7640	344.752	-18.511	0.000	-7058.306	-5705.222
dayofweek_Thurs	387.7770	336.822	1.151	0.250	-273.202	1048.757
dayofweek_Tues	-401.5068	343.769	-1.168	0.243	-1076.118	273.105
dayofweek_Wed	606.1591	341.981	1.772	0.077	-64.944	1277.263
holiday_False	-985.5413	468.510	-2.104	0.036	-1904.946	-66.137
holiday_True	-1.157e+04	624.110	-18.534	0.000	-1.28e+04	-1.03e+04
month_Apr	-2258.5999	524.291	-4.308	0.000	-3287.468	-1229.732
month_Aug	2209.4244	579.423	3.813	0.000	1072.365	3346.484
month_Dec	-5023.3082	482.377	-10.414	0.000	-5969.925	-4076.691
month_Feb	-7612.2728	699.954	-10.875	0.000	-8985.863	-6238.682
month_Jan	-7134.2280	606.453	-11.764	0.000	-8324.331	-5944.125
month_Jul	461.3105	617.169	0.747	0.455	-749.821	1672.442
month_Jun	2268.7701	612.624	3.703	0.000	1066.557	3470.983
month_Mar	-5783.7715	508.398	-11.376	0.000	-6781.451	-4786.091
month_May	1488.4128	565.880	2.630	0.009	377.930	2598.895
month_Nov	-188.1675	454.469	-0.414	0.679	-1080.017	703.682
month_Oct	4373.2223	434.020	10.076	0.000	3521.501	5224.944
month_Sep	4646.5193	512.235	9.071	0.000	3641.310	5651.728
n_stations	68.9045	2.888	23.855	0.000	63.236	74.573

```
=====
```

Omnibus:	56.833	Durbin-Watson:	1.150
Prob(Omnibus):	0.000	Jarque-Bera (JB):	123.489
Skew:	-0.345	Prob(JB):	1.53e-27
Kurtosis:	4.577	Cond. No.	1.06e+17

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.68e-26. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

MSE: 17254735.842562452

2.2.3 Backward selection

2.2.3.1 Useful functions

In [329]: # useful functions

```
def whichColumnShouldIDrop(modelResult):
    pSeries = modelResult.pvalues
    print(f"There were {len(pSeries) - 1} variables in the model.")
    highestPvalue = max(pSeries)
    varName = pSeries[pSeries == highestPvalue].keys()[0]
    print(f"The one with the highest pvalue of '{highestPvalue}' was '{varName}'")
    print(f"The column to drop should therefore be {varName}")
```

In [330]:

```
def allColumnsWithXInName(x):
    allCols = list(train.columns)
    return list(filter(lambda col: x in col, allCols))
```

In [331]:

```
def calculateTestMSE(modelResult, y_test, X_test):
    y_pred = modelResult.predict(X_test)
    subs = [(y_test.trips[i] - y_pred[i]) ** 2 for i in range(len(y_test))]
    p = len(modelResult.pvalues) - 1
    return sum(subs) / (len(y_test) - p)
```

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
 - ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

2.2.3.2 Best model $p = 5$

Five rounds of selection

```
In [332]: whichColumnShouldIDrop(allCovariatesResult)
```

There were 28 variables in the model.
The one with the highest pvalue of '0.989853175886976' was 'SNOW'
The column to drop should therefore be SNOW

```
In [333]: covariates = train.drop(columns=["trips", "SNOW"])
X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
firstSelectionResult = model.fit()
print(firstSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.873
Model:	OLS	Adj. R-squared:	0.870
Method:	Least Squares	F-statistic:	280.0
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	0.00
Time:	22:08:46	Log-Likelihood:	-9747.3
No. Observations:	1001	AIC:	1.954e+04
Df Residuals:	976	BIC:	1.967e+04
Df Model:	24		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.255e+04	750.890	-16.716	0.000	-1.4e+04	-1.11e+04
AWND	0.7647	0.244	3.140	0.002	0.287	1.243
PRCP	-8213.1439	384.700	-21.349	0.000	-8968.079	-7458.209
SNWD	-215.0749	62.326	-3.451	0.001	-337.383	-92.767
TMAX	352.8805	29.693	11.884	0.000	294.610	411.151
TMIN	-73.3368	32.773	-2.238	0.025	-137.650	-9.024
dayofweek_Fri	-309.3256	336.822	-0.918	0.359	-970.304	351.653
dayofweek_Mon	-1038.9093	341.093	-3.046	0.002	-1708.269	-369.549
dayofweek_Sat	-5414.6308	341.379	-15.861	0.000	-6084.553	-4744.709
dayofweek_Sun	-6381.7238	344.561	-18.521	0.000	-7057.890	-5705.558
dayofweek_Thurs	387.9528	336.366	1.153	0.249	-272.131	1048.037
dayofweek_Tues	-401.2177	342.841	-1.170	0.242	-1074.008	271.573
dayofweek_Wed	606.0606	341.718	1.774	0.076	-64.526	1276.647
holiday_False	-984.9759	466.158	-2.113	0.035	-1899.764	-70.188
holiday_True	-1.157e+04	623.253	-18.559	0.000	-1.28e+04	-1.03e+04
month_Apr	-2258.8742	523.579	-4.314	0.000	-3286.343	-1231.405
month_Aug	2209.6592	578.832	3.817	0.000	1073.761	3345.558
month_Dec	-5023.3601	482.113	-10.419	0.000	-5969.457	-4077.263
month_Feb	-7612.0582	699.392	-10.884	0.000	-8984.544	-6239.572
month_Jan	-7133.8050	605.230	-11.787	0.000	-8321.507	-5946.103
month_Jul	461.5060	616.661	0.748	0.454	-748.628	1671.640
month_Jun	2268.8147	612.300	3.705	0.000	1067.238	3470.391
month_Mar	-5783.7481	508.134	-11.382	0.000	-6780.909	-4786.587
month_May	1488.4185	565.590	2.632	0.009	378.507	2598.330
month_Nov	-188.3026	454.112	-0.415	0.678	-1079.450	702.845
month_Oct	4373.2353	433.797	10.081	0.000	3521.954	5224.517
month_Sep	4646.7208	511.727	9.080	0.000	3642.508	5650.933
n_stations	68.9022	2.881	23.914	0.000	63.248	74.556

```
=====
```

Omnibus:	56.836	Durbin-Watson:	1.150
Prob(Omnibus):	0.000	Jarque-Bera (JB):	123.434
Skew:	-0.345	Prob(JB):	1.57e-27
Kurtosis:	4.576	Cond. No.	1.06e+17

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.68e-26. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [334]: whichColumnShouldIDrop(firstSelectionResult)
```

There were 27 variables in the model.
The one with the highest pvalue of '0.6784808069754316' was 'month_Nov'
The column to drop should therefore be month_Nov

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regressi
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and
 - ▼ 2.1 Problem 2.a: Read in and merge
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regressor
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of column:
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find me
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Ga
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and impo
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator fu
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [335]: # remove all the month variables
covariates = train.drop(columns=["trips", "SNOW"] + allColumnsWithName("month"))
X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
secondSelectionResult = model.fit()
print(secondSelectionResult.summary())
```

```
OLS Regression Results
=====
Dep. Variable:      trips      R-squared:      0.820
Model:              OLS      Adj. R-squared:    0.817
Method:             Least Squares      F-statistic:    345.3
Date:               Tue, 17 Sep 2019      Prob (F-statistic): 0.00
Time:               22:08:46      Log-Likelihood: -9923.2
No. Observations:   1001      AIC:          1.987e+04
Df Residuals:       987      BIC:          1.994e+04
Df Model:           13
Covariance Type:    nonrobust
=====
                    coef      std err      t      P>|t|      [0.025      0.975]
-----
const             -1.674e+04    846.118    -19.788    0.000    -1.84e+04    -1.51e+04
AWND                0.6034         0.286     2.106    0.035         0.041         1.166
PRCP             -8684.1951    453.859    -19.134    0.000    -9574.835    -7793.555
SNWD             -411.2585         61.553     -6.681    0.000     -532.048    -290.469
TMAX              408.0427         32.228     12.661    0.000       344.799       471.286
TMIN               37.2733         34.333     1.086    0.278       -30.100      104.646
dayofweek_Fri     -770.4382        397.535     -1.938    0.053    -1550.548         9.672
dayofweek_Mon    -1664.1143        402.489     -4.135    0.000    -2453.948     -874.281
dayofweek_Sat    -5911.2112        403.187    -14.661    0.000    -6702.413    -5120.010
dayofweek_Sun    -7139.6523        404.772    -17.639    0.000    -7933.965    -6345.340
dayofweek_Thurs  -164.0064         396.880     -0.413    0.680     -942.832       614.819
dayofweek_Tues   -1060.1334        403.953     -2.624    0.009    -1852.839    -267.428
dayofweek_Wed    -33.6629         402.266     -0.084    0.933     -823.057       755.732
holiday_False    -2812.0451        537.708     -5.230    0.000    -3867.227    -1756.863
holiday_True     -1.393e+04        717.202    -19.424    0.000    -1.53e+04    -1.25e+04
n_stations         60.2811         2.916     20.672    0.000         54.559         66.003
=====
Omnibus:           23.647      Durbin-Watson:      0.878
Prob(Omnibus):     0.000      Jarque-Bera (JB):    44.622
Skew:              -0.120      Prob(JB):            2.04e-10
Kurtosis:          4.006      Cond. No.            6.76e+18
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The smallest eigenvalue is 6.57e-30. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
```

```
In [336]: whichColumnShouldIDrop(secondSelectionResult)
```

There were 15 variables in the model.
The one with the highest pvalue of '0.9333252469205151' was 'dayofweek_Wed'
The column to drop should therefore be dayofweek_Wed

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs SNOW
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression model
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [337]: # remove all the day variables
covariates = train.drop(
    columns=["trips", "SNOW"] +
    allColumnsWithXInName("month") +
    allColumnsWithXInName("dayofweek"))

X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
thirdSelectionResult = model.fit()
print(thirdSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.767
Model:	OLS	Adj. R-squared:	0.765
Method:	Least Squares	F-statistic:	466.3
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	1.21e-308
Time:	22:08:46	Log-Likelihood:	-10052.
No. Observations:	1001	AIC:	2.012e+04
Df Residuals:	993	BIC:	2.016e+04
Df Model:	7		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.808e+04	1049.460	-17.229	0.000	-2.01e+04	-1.6e+04
AWND	1.0357	0.323	3.202	0.001	0.401	1.671
PRCP	-8328.7629	512.770	-16.243	0.000	-9335.001	-7322.525
SNWD	-406.4006	69.802	-5.822	0.000	-543.376	-269.425
TMAX	395.5854	36.289	10.901	0.000	324.373	466.797
TMIN	50.9168	38.676	1.316	0.188	-24.980	126.813
holiday_False	-4239.3928	631.843	-6.710	0.000	-5479.294	-2999.491
holiday_True	-1.384e+04	842.140	-16.436	0.000	-1.55e+04	-1.22e+04
n_stations	61.4942	3.306	18.599	0.000	55.006	67.982

```
=====
```

Omnibus:	4.582	Durbin-Watson:	0.921
Prob(Omnibus):	0.101	Jarque-Bera (JB):	5.352
Skew:	-0.048	Prob(JB):	0.0689
Kurtosis:	3.345	Cond. No.	1.10e+18

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.48e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [338]: whichColumnShouldIDrop(thirdSelectionResult)
```

There were 8 variables in the model.
 The one with the highest pvalue of '0.18831446195624318' was 'TMIN'
 The column to drop should therefore be TMIN

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression model
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model p == 5
 - 2.2.3.3 Best model p == 4
 - 2.2.3.4 Best model p == 3
 - 2.2.3.5 Best model p == 2
 - 2.2.3.6 Best model p == 1
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [339]: covariates = train.drop(
           columns=["trips", "SNOW", "TMIN"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek"))

X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
fourthSelectionResult = model.fit()
print(fourthSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.766
Model:	OLS	Adj. R-squared:	0.765
Method:	Least Squares	F-statistic:	543.3
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	1.18e-309
Time:	22:08:46	Log-Likelihood:	-10053.
No. Observations:	1001	AIC:	2.012e+04
Df Residuals:	994	BIC:	2.015e+04
Df Model:	6		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.83e+04	1035.981	-17.669	0.000	-2.03e+04	-1.63e+04
AWND	1.0467	0.324	3.235	0.001	0.412	1.682
PRCP	-8280.9945	511.674	-16.184	0.000	-9285.079	-7276.910
SNWD	-417.5572	69.311	-6.024	0.000	-553.569	-281.545
TMAX	441.3231	10.486	42.088	0.000	420.747	461.900
holiday_False	-4374.5435	623.678	-7.014	0.000	-5598.419	-3150.668
holiday_True	-1.393e+04	839.753	-16.589	0.000	-1.56e+04	-1.23e+04
n_stations	61.3480	3.306	18.559	0.000	54.861	67.835

```
=====
```

Omnibus:	3.885	Durbin-Watson:	0.925
Prob(Omnibus):	0.143	Jarque-Bera (JB):	4.422
Skew:	-0.037	Prob(JB):	0.110
Kurtosis:	3.317	Cond. No.	7.95e+16

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 4.76e-26. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [340]: whichColumnShouldIDrop(fourthSelectionResult)
```

There were 7 variables in the model.
The one with the highest pvalue of '0.0012545608107361673' was 'AWND'
The column to drop should therefore be AWND

```
In [341]: covariates = train.drop(
           columns=["trips", "SNOW", "TMIN", "AWND"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek"))

X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
fifthSelectionResult = model.fit()
print(fifthSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.764
Model:	OLS	Adj. R-squared:	0.763
Method:	Least Squares	F-statistic:	643.7
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	8.01e-309
Time:	22:08:46	Log-Likelihood:	-10058.
No. Observations:	1001	AIC:	2.013e+04
Df Residuals:	995	BIC:	2.016e+04
Df Model:	5		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.822e+04	1040.561	-17.509	0.000	-2.03e+04	-1.62e+04
PRCP	-8252.3813	514.026	-16.054	0.000	-9261.080	-7243.683
SNWD	-417.8816	69.640	-6.001	0.000	-554.539	-281.224
TMAX	442.2684	10.531	41.995	0.000	421.602	462.935
holiday_False	-4354.2871	626.606	-6.949	0.000	-5583.909	-3124.666
holiday_True	-1.387e+04	843.497	-16.438	0.000	-1.55e+04	-1.22e+04
n_stations	60.7974	3.317	18.329	0.000	54.288	67.306

```
=====
```

Omnibus:	4.606	Durbin-Watson:	0.924
Prob(Omnibus):	0.100	Jarque-Bera (JB):	5.349
Skew:	-0.053	Prob(JB):	0.0690
Kurtosis:	3.342	Cond. No.	7.58e+17

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Contents


- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMIN
 - 2.1.3.2 Pair plot of all numeric
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of column:
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p == 5$
 - 2.2.3.3 Best model $p == 4$
 - 2.2.3.4 Best model $p == 3$
 - 2.2.3.5 Best model $p == 2$
 - 2.2.3.6 Best model $p == 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find me
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standard
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and impute
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

At the end of five rounds of selection, the model with the following variables was chosen: [PRCP, SNWD, TMAX, holiday, n_stations]

Pair plot

```
In [342]: sns.pairplot(train.drop(
    columns=["SNOW", "TMIN", "AWND"] +
    allColumnsWithName("month") +
    allColumnsWithName("dayofweek")),
    y_vars=["trips"],
    x_vars=["PRCP", "SNWD", "TMAX", "holiday_False", "holiday_True", "n_stations"])
```

Out[342]: <seaborn.axisgrid.PairGrid at 0x12c072150>



Predictions

R^2 and MSE

```
In [343]: fifthSelectionResult.mse_resid
```

Out[343]: 31482485.09965653

```
In [344]: X_test = train.drop(
    columns=["SNOW", "TMIN", "AWND"] +
    allColumnsWithName("month") +
    allColumnsWithName("dayofweek"))

y_test = train[["trips"]]

calculateTestMSE(fifthSelectionResult, y_test, X_test)
```

Out[344]: 12795186806733.623

Training data

R^2 according to the model summary was: 0.764

MSE was: 31482485.09965653

Test data

R^2 : According to the rubric and Parker's answer on Piazza, this is a meaningless statistic so I have not computed it.

MSE was: 12795186806733.623

2.2.3.3 Best model $p == 4$

Another round of selection

```
In [345]: whichColumnShouldIDrop(fifthSelectionResult)
```

There were 6 variables in the model.
The one with the highest pvalue of '2.749583091977293e-09' was 'SNWD'
The column to drop should therefore be SNWD

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [346]: covariates = train.drop(
          columns=["trips", "SNOW", "TMIN", "AWND", "SNWD"] +
          allColumnsWithXInName("month") +
          allColumnsWithXInName("dayofweek"))

X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
sixthSelectionResult = model.fit()
print(sixthSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.755
Model:	OLS	Adj. R-squared:	0.754
Method:	Least Squares	F-statistic:	768.6
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	1.28e-302
Time:	22:08:47	Log-Likelihood:	-10076.
No. Observations:	1001	AIC:	2.016e+04
Df Residuals:	996	BIC:	2.019e+04
Df Model:	4		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-2.065e+04	975.130	-21.178	0.000	-2.26e+04	-1.87e+04
PRCP	-8236.0150	522.974	-15.748	0.000	-9262.272	-7209.758
TMAX	471.0558	9.539	49.385	0.000	452.338	489.774
holiday_False	-5666.5170	597.439	-9.485	0.000	-6838.900	-4494.134
holiday_True	-1.498e+04	836.951	-17.904	0.000	-1.66e+04	-1.33e+04
n_stations	65.2022	3.291	19.812	0.000	58.744	71.660

```
=====
```

Omnibus:	3.845	Durbin-Watson:	0.916
Prob(Omnibus):	0.146	Jarque-Bera (JB):	4.236
Skew:	-0.058	Prob(JB):	0.120
Kurtosis:	3.297	Cond. No.	8.09e+17

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 2.02e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

R^2 and MSE

```
In [347]: sixthSelectionResult.mse_resid
```

```
Out[347]: 32589036.105076745
```

```
In [348]: X_test = train.drop(
          columns=["SNOW", "TMIN", "AWND", "SNWD"] +
          allColumnsWithXInName("month") +
          allColumnsWithXInName("dayofweek"))

y_test = train["trips"]

calculateTestMSE(sixthSelectionResult, y_test, X_test)
```

```
Out[348]: 18657357119633.902
```

Training data

R^2 according to the model summary was: 0.755

MSE was: 32589036.105076745

Test data

MSE was: 18657357119633.902

2.2.3.4 Best model $p = 3$

Another selection

```
In [349]: whichColumnShouldIDrop(sixthSelectionResult)
```

There were 5 variables in the model.
The one with the highest pvalue of '1.7255263964426543e-20' was 'holiday_False'
The column to drop should therefore be holiday_False

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs PRCP
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression model
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [350]: covariates = train.drop(
           columns=["trips", "SNOW", "TMIN", "AWND", "SNWD"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek") +
           allColumnsWithXInName("holiday"))

X = sm.add_constant(covariates)
y = train["trips"]
model = sm.OLS(y, X)
seventhSelectionResult = model.fit()
print(seventhSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
```

Dep. Variable:	trips	R-squared:	0.737
Model:	OLS	Adj. R-squared:	0.736
Method:	Least Squares	F-statistic:	931.2
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	1.57e-288
Time:	22:08:47	Log-Likelihood:	-10112.
No. Observations:	1001	AIC:	2.023e+04
Df Residuals:	997	BIC:	2.025e+04
Df Model:	3		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-2.685e+04	1432.620	-18.743	0.000	-2.97e+04	-2.4e+04
PRCP	-8085.1685	541.629	-14.927	0.000	-9148.033	-7022.304
TMAX	476.5096	9.863	48.315	0.000	457.156	495.863
n_stations	64.9314	3.410	19.040	0.000	58.239	71.623

```
=====
```

Omnibus:	10.038	Durbin-Watson:	0.832
Prob(Omnibus):	0.007	Jarque-Bera (JB):	11.724
Skew:	-0.157	Prob(JB):	0.00285
Kurtosis:	3.427	Cond. No.	2.79e+03

```
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.79e+03. This might indicate that there are strong multicollinearity or other numerical problems.

R^2 and MSE

```
In [351]: seventhSelectionResult.mse_resid
```

```
Out[351]: 34994570.75404534
```

```
In [352]: X_test = train.drop(
           columns=["SNOW", "TMIN", "AWND", "SNWD"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek") +
           allColumnsWithXInName("holiday"))

y_test = train["trips"]

calculateTestMSE(seventhSelectionResult, y_test, X_test)
```

```
Out[352]: 1998366500725.9263
```

Training data

R^2 according to the model summary was: 0.737

MSE was: 34994570.75404534

Test data

MSE was: 1998366500725.9263

2.2.3.5 Best model $p = 2$

Another selection

```
In [353]: whichColumnShouldIDrop(seventhSelectionResult)
```

There were 3 variables in the model.
 The one with the highest pvalue of '1.2561082989927261e-45' was 'PRCP'
 The column to drop should therefore be PRCP

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression model
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [354]: covariates = train.drop(
           columns=["trips", "SNOW", "TMIN", "AWND", "SNWD", "PRCP"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek") +
           allColumnsWithXInName("holiday"))

X = sm.add_constant(covariates)
y = train[["trips"]]
model = sm.OLS(y, X)
eighthSelectionResult = model.fit()
print(eighthSelectionResult.summary())
```

OLS Regression Results

Dep. Variable:	trips	R-squared:	0.678			
Model:	OLS	Adj. R-squared:	0.678			
Method:	Least Squares	F-statistic:	1052.			
Date:	Tue, 17 Sep 2019	Prob (F-statistic):	1.93e-246			
Time:	22:08:47	Log-Likelihood:	-10213.			
No. Observations:	1001	AIC:	2.043e+04			
Df Residuals:	998	BIC:	2.045e+04			
Df Model:	2					
Covariance Type:	nonrobust					

	coef	std err	t	P> t	[0.025	0.975]
const	-2.874e+04	1577.678	-18.215	0.000	-3.18e+04	-2.56e+04
TMAX	479.3314	10.902	43.968	0.000	457.938	500.724
n_stations	67.0363	3.767	17.796	0.000	59.644	74.428

Omnibus: 18.280 Durbin-Watson: 0.958
 Prob(Omnibus): 0.000 Jarque-Bera (JB): 18.879
 Skew: -0.335 Prob(JB): 7.95e-05
 Kurtosis: 3.054 Cond. No. 2.78e+03

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 2.78e+03. This might indicate that there are strong multicollinearity or other numerical problems.

R^2 and MSE

```
In [355]: eighthSelectionResult.mse_resid
```

Out[355]: 42772978.27762087


```
In [356]: X_test = train.drop(
           columns=["SNOW", "TMIN", "AWND", "SNWD", "PRCP"] +
           allColumnsWithXInName("month") +
           allColumnsWithXInName("dayofweek") +
           allColumnsWithXInName("holiday"))

y_test = train[["trips"]]

calculateTestMSE(eighthSelectionResult, y_test, X_test)
```

Out[356]: 260141784482.93845

Training data

R^2 according to the model summary was: 0.678

MSE was: 42772978.27762087

Test data

MSE was: 260141784482.93845

2.2.3.6 Best model $p = 1$

Another selection

```
In [357]: whichColumnShouldIDrop(eighthSelectionResult)
```

There were 2 variables in the model.
 The one with the highest pvalue of '9.615354198374857e-62' was 'n_stations'
 The column to drop should therefore be n_stations

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression model
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find median
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a.
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b.
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [358]: covariates = train.drop(
           columns=["trips", "SNOW", "TMIN", "AWND", "SNWD", "PRCP", "n_stations"] +
           allColumnsWithName("month") +
           allColumnsWithName("dayofweek") +
           allColumnsWithName("holiday"))

X = sm.add_constant(covariates)
y = train[["trips"]]
model = sm.OLS(y, X)
ninthSelectionResult = model.fit()
print(ninthSelectionResult.summary())
```

```
=====
                        OLS Regression Results
=====
Dep. Variable:          trips          R-squared:                0.576
Model:                  OLS          Adj. R-squared:            0.576
Method:                 Least Squares   F-statistic:              1358.
Date:                   Tue, 17 Sep 2019   Prob (F-statistic):       2.22e-188
Time:                   22:08:48         Log-Likelihood:           -10351.
No. Observations:       1001            AIC:                     2.071e+04
Df Residuals:           999            BIC:                     2.072e+04
Df Model:                1
Covariance Type:        nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const      -3659.5444      813.732      -4.497      0.000     -5256.364     -2062.725
TMAX         458.0092       12.430       36.846      0.000       433.616       482.402
=====
Omnibus:            12.667      Durbin-Watson:           0.735
Prob(Omnibus):       0.002      Jarque-Bera (JB):        12.782
Skew:                0.263      Prob(JB):                0.00168
Kurtosis:            3.170      Cond. No.                 225.
=====
```

```
Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

R^2 and MSE

```
In [359]: ninthSelectionResult.mse_resid
```

```
Out[359]: 56289579.39584944
```

```
In [360]: X_test = train.drop(
           columns=["SNOW", "TMIN", "AWND", "SNWD", "PRCP", "n_stations"] +
           allColumnsWithName("month") +
           allColumnsWithName("dayofweek") +
           allColumnsWithName("holiday"))

y_test = train[["trips"]]

calculateTestMSE(ninthSelectionResult, y_test, X_test)
```

```
Out[360]: 152850992425903.4
```

Training data

R^2 according to the model summary was: 0.576

MSE was: 56289579.39584944

Test data

MSE was: 152850992425903.4

2.3 Problem 2.c: KNN Classification

Now we will transform the outcome variable to allow us to do classification. Create a new vector Y with entries:

$$Y[i] = \mathbf{1}\{\text{trips}[i] > \text{median}(\text{trips})\}$$

Use the median of the variable from the full data (training and test combined). After computing the binary outcome variable Y , you should drop the original trips variable from the data.

2.3.1 Compute binary Y, find median, drop "trips"

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find mean
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardization
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [361]: median = np.median(np.concatenate((test.trips, train.trips)))

knnYTrain = [1 if trips > median else 0 for trips in train.trips]
knnXTrain = train.drop(columns=["trips"])

knnYTest = [1 if trips > median else 0 for trips in test.trips]
knnXTest = test.drop(columns=["trips"])
```

2.3.2 Drop categorical variables

Recall that in k -nearest neighbors classification, the predicted value \hat{Y} of X is the majority vote of the labels for the k nearest neighbors X_i to X . We will use the Euclidean distance as our measure of distance between points. Note that the Euclidean distance doesn't make much sense for factor variables, so just drop the predictors that are categorical for this problem. Standardize the numeric predictors so that they have mean zero and constant standard deviation.

```
In [362]: knnXTrain = knnXTrain.drop(columns=(
    allColumnsWithXInName("month") +
    allColumnsWithXInName("dayofweek") +
    allColumnsWithXInName("holiday")))

knnXTest = knnXTest.drop(columns=(
    allColumnsWithXInName("month") +
    allColumnsWithXInName("dayofweek") +
    allColumnsWithXInName("holiday")))
```

2.3.3 Feature scaling: Standardization

```
In [363]: from sklearn.preprocessing import StandardScaler
scX = StandardScaler()
```

$$X = \frac{X_{test} - \mu_{train}}{\sigma_{train}}$$

The reason we just transform test data instead of `fit_transform` is so that we can have the same train parameters apply to the test data (the `scX` maintains some internal state).

```
In [364]: # why do fit_transform on train and just transform on test?
knnXTrain = scX.fit_transform(knnXTrain)
knnXTest = scX.transform(knnXTest)
```

2.3.4 Run the KNN

You may use the `KNeighborsClassifier` function from the `sklearn.neighbors` module to perform k -nearest neighbor classification, using as the neighbors the labeled points in the training set. Fit a classifier for $k = 1 : 50$, and find the mis-classification rate on both the training and test sets for each k . On a single plot, show the training set error and the test set error as a function of k . How would you choose the optimal k ? Comment on your findings, and in particular on the possibility of overfitting.

```
In [365]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
```


```
In [366]: neighbors_settings = range(1, 51)
training_accuracy = []
test_accuracy = []
for n_neighbors in neighbors_settings:
    # instantiate a classifier
    classifier = KNeighborsClassifier(n_neighbors=n_neighbors)
    # fit the classifier to the training data
    classifier.fit(knnXTrain, knnYTrain)
    # find the training set accuracy to the actual labels (did it learn the model well?)
    # and find the test set accuracy to the actual labels (did it generalize well?)
    training_accuracy.append(1 - classifier.score(knnXTrain, knnYTrain))
    test_accuracy.append(1 - classifier.score(knnXTest, knnYTest))
```

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian Mixture
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [367]: plt.plot(neighbors_settings, training_accuracy, label="training error")
plt.plot(neighbors_settings, test_accuracy, label="test error")
plt.ylabel("Error")
plt.xlabel("n_neighbors")
plt.legend()
```

Out[367]: <matplotlib.legend.Legend at 0x11bf26f50>



2.3.5 Best model

The best model is the one that minimizes the test error/misclassification, around $k = 40$.

We might get over-fitting if we just try to minimize training misclassification, because that corresponds to $k = 1$.

3 Problem 3: Classification for a Gaussian Mixture (25 points)

A Gaussian mixture model is a random combination of multiple Gaussians. Specifically, we can generate n data points from such a distribution in the following way. First generate labels Y_1, \dots, Y_n according to

$$Y_i = \begin{cases} 0 & \text{with probability } 1/2 \\ 1 & \text{with probability } 1/2. \end{cases}$$

Then, generate the data X_1, \dots, X_n according to

$$X_i \sim \begin{cases} N(\mu_0, \sigma_0^2) & \text{if } Y_i = 0 \\ N(\mu_1, \sigma_1^2) & \text{if } Y_i = 1. \end{cases}$$

Given such data $\{X_i\}$, we may wish to recover the true labels Y_i , which is a classification task.

3.1 Problem 3.a.

Suppose the parameters of the above model are: $\mu_0 = 0, \mu_1 = 3, \sigma_0^2 = \sigma_1^2 = 1$. Then the Bayes classifier is given by

$$f(X) = I\{X > 1.5\},$$

where I is the indicator function (take note of the 1.5, and it's relation with the means of the two Normal distributions).

Now generate $n = 2000$ data points from this dataset. Plot a histogram of the X 's. This histogram is meant to be a sanity check for you; it should help you verify that you've generated the data properly.

3.1.1 Useful functions and imports

```
In [368]: import random
from sklearn.model_selection import train_test_split
```

3.1.2 Generate Ys

```
In [369]: gaussian_y = [random.choice([0, 1]) for x in range(0, 2000)]
```

3.1.3 Generate Xs


```
In [370]: gaussian_x = [np.random.normal(loc=0, scale=1)
                        if gaussian_y[x] == 0 else
                        np.random.normal(loc=3, scale=1)
                        for x in range(0, 2000)]
```

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator function
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

```
In [371]: # sanity check
plt.hist(gaussian_x)
```

```
Out[371]: (array([ 20.,  75., 248., 374., 283., 234., 327., 280., 127.,  32.]),
array([-3.0577953, -2.17814123, -1.29848716, -0.4188331,  0.46082097,
        1.34047504,  2.2201291,  3.09978317,  3.97943724,  4.8590913,
        5.73874537])),
<a list of 10 Patch objects>)
```



3.1.4 Train-test split

Set aside a randomly-selected test set of $n/5$ points.

```
In [372]: X_train, X_test, y_train, y_test = train_test_split(gaussian_x, gaussian_y, test_size=0.2)
```

3.1.5 Calculate group means

We will refer to the rest of the data as the training data. Use the labels of the training data to calculate the group means. That is, calculate the mean value of all the X_i 's in the training data with label $Y_i = 0$. Call this sample mean $\hat{\mu}_0$. Do the same thing to find $\hat{\mu}_1$. To be explicit, let $C_j = \{i : Y_i = j\}$, and define

$$\hat{\mu}_j = \frac{1}{|C_j|} \sum_{i \in C_j} X_i$$

```
In [373]: def calculateMeanOfAllXWithLabel(label, X_train, y_train):
X_train_with_label = []
for index, y in enumerate(y_train):
    if y == label:
        X_train_with_label.append(X_train[index])
return np.mean(X_train_with_label)
```

```
In [374]: # calculate u0 and u1
u_0_hat = calculateMeanOfAllXWithLabel(0, X_train, y_train)
u_1_hat = calculateMeanOfAllXWithLabel(1, X_train, y_train)
```

3.1.6 Classify using indicator function

Now classify the data in your test set. To do this, recall that your rule in Part a. depended on the true data means $\mu_0 = 0$ and $\mu_1 = 3$. Plug in the sample means $\hat{\mu}_j$ instead. Evaluate the estimator's performance using the loss:

$$\frac{1}{n} \sum_{i=1}^n 1\{\hat{Y}_i \neq Y_i\}$$

```
In [375]: def indicatorFunction(x, mean1, mean2):
    if x > ((mean1 + mean2) / 2):
        return 1
    return 0
```

```
In [376]: y_pred = [indicatorFunction(x, u_0_hat, u_1_hat) for x in X_test]
```

3.1.7 Compute error rate

```
In [377]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
```

```
In [378]: misclassified_count = cm[1][0] + cm[0][1]
error_rate = misclassified_count / len(y_pred)
print(f"Error rate is: {error_rate}")
```

Error rate is: 0.0875

Contents

- ▼ 0.1 S&DS 355/555: Assignment 1
 - 0.1.1 NetID: sa857
 - 0.1.2 Imports and settings
- ▼ 1 Problem 1: Simple Linear Regression
 - ▼ 1.1 Problem 1.a:
 - 1.1.1 Answer
 - ▼ 1.2 Problem 1.b:
 - 1.2.1 Answer
 - ▼ 1.3 Problem 1.c:
 - 1.3.1 Answer
 - ▼ 1.4 Problem 1.d:
 - 1.4.1 Answer
- ▼ 2 Problem 2: Linear regression and model selection
 - ▼ 2.1 Problem 2.a: Read in and merge data
 - 2.1.1 Import data
 - 2.1.2 Merge data
 - ▼ 2.1.3 Explore data
 - 2.1.3.1 Scatterplot of TMIN vs TMAX
 - 2.1.3.2 Pair plot of all numeric variables
 - ▼ 2.2 Problem 2.b: Linear regression
 - ▼ 2.2.1 Preprocessing
 - 2.2.1.1 Categorical variables
 - 2.2.1.2 Reordering of columns
 - 2.2.2 Model all covariates
 - ▼ 2.2.3 Backward selection
 - 2.2.3.1 Useful functions
 - 2.2.3.2 Best model $p = 5$
 - 2.2.3.3 Best model $p = 4$
 - 2.2.3.4 Best model $p = 3$
 - 2.2.3.5 Best model $p = 2$
 - 2.2.3.6 Best model $p = 1$
 - ▼ 2.3 Problem 2.c: KNN Classification
 - 2.3.1 Compute binary Y, find means
 - 2.3.2 Drop categorical variables
 - 2.3.3 Feature scaling: Standardize
 - 2.3.4 Run the KNN
 - 2.3.5 Best model
- ▼ 3 Problem 3: Classification for a Gaussian Process
 - ▼ 3.1 Problem 3.a:
 - 3.1.1 Useful functions and imports
 - 3.1.2 Generate Ys
 - 3.1.3 Generate Xs
 - 3.1.4 Train-test split
 - 3.1.5 Calculate group means
 - 3.1.6 Classify using indicator functions
 - 3.1.7 Compute error rate
 - ▼ 3.2 Problem 3.b:
 - 3.2.1 Plot error rates
 - 3.2.2 Comments

3.2 Problem 3.b.

Now you train and evaluate classifiers for training sets of increasing size n , as specified below. For each n , you should

1. Generate a training set of size n from the above model (with the same parameters).
2. Generate a test set of size 10,000. Note that the test set itself will change on each round, but the size will always be the same: 10,000.
3. Compute the sample means on the training data.
4. Classify the test data as described in Part c.
5. Compute the error rate.

Plot the error rate as a function of n . Comment on your findings. What is happening to the error rate as n grows?


```
In [379]: ## -- please write your answer here. -- ##
seq_n = np.arange(start = 2000, stop = 20000, step = 20)
error_rates = []

In [380]: for n_count in seq_n:
# create labels according to n_count + 10000 == total count
gaussian_y = [random.choice([0, 1]) for x in range(0, n_count + 10000)]
# create feature data
gaussian_x = [np.random.normal(loc=0, scale=1)
               if gaussian_y[x] == 0 else
               np.random.normal(loc=3, scale=1)
               for x in range(0, n_count + 10000)]
# train test split
X_train, X_test, y_train, y_test = train_test_split(gaussian_x, gaussian_y, test_size=10000)
# calculate means
u_0_hat = calculateMeanOfAllXWithLabel(0, X_train, y_train)
u_1_hat = calculateMeanOfAllXWithLabel(1, X_train, y_train)
# make the predictions
y_pred = [indicatorFunction(x, u_0_hat, u_1_hat) for x in X_test]
# compute error rate
cm = confusion_matrix(y_test, y_pred)
misclassified_count = cm[1][0] + cm[0][1]
error_rate = misclassified_count / len(y_pred)
error_rates.append(error_rate)
```

3.2.1 Plot error rates


```
In [381]: plt.plot(seq_n, error_rates)

Out[381]: <matplotlib.lines.Line2D at 0x12ff76d50>
```



```
In [382]: plt.hist(error_rates)

Out[382]: (array([ 3., 17., 47., 111., 178., 251., 164., 90., 29., 10.]),
array([0.0586, 0.06014, 0.06168, 0.06322, 0.06476, 0.0663, 0.06784,
0.06938, 0.07092, 0.07246, 0.074 ]),
<a list of 10 Patch objects>)
```



3.2.2 Comments

Error rate is fairly constant, and normally distributed around 0.067.