Contents 2 4

```
1 Assignment 2
```

- ▼ 2 Problem 1: Spam, wonderful spam! 2.0.1 A Note about Features 2.0.2 Missing Values ▼ 2.1 Part a
- 2.1.1 Scaling
- 2.1.2 Fit KNR to non-NaN frames
- ▼ 2.3 Part c
 - 2.3.1 KNN all predictors except ca 2.3.2 KNN - all predictors
 - 2.3.3 Logistic Regression all predic
- ▼ 2.3.4 Logistic Regression all predic 2.3.4.1 Significant predictors 2.3.5 Export to CSV
- ▼ 3 Problem 2: Gradient Descent
 - 3.1 Part a
 - 3.2 Part b
 - 3.3 Part c 3.4 Part d
- ▼ 4 Problem 3: Cross-Validation
 - 4.1 Part a
 - 4.2 Part b
 - 4.3 Part c
 - 4.4 Part d
 - 4.5 Part e 4.6 Part f

In [240]: import matplotlib.pyplot as plt import numpy as np import pandas as pd import math import matn
from perspective import psp
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier from sklearn.linear_model import LogisticRegression

In [176]: %load_ext line_profiler The line_profiler extension is already loaded. To reload it, use: %reload_ext line_profiler

Notice: Unfortunately, due to sickness and multiple commitments, I fell behind on my PSET this week,

I have completed Problem 1 and parts of Problem 3. Thank you for your understanding.

1 Assignment 2

Due: 26th Sept, 11:59pm

2 Problem 1: Spam, wonderful spam!

The dataset consists of a collection of 57 features relating to about 4600 emails and a label of whether or not the email is considered spam. You have a training set containing about 70% of the data and a test set containing about 30% of the data. Your job is to build effective spam classification rules using the predictors

2.0.1 A Note about Features

The column names (in the first row of each .csv file) are fairly self-explanatory

- Some variables are named word_freq_(word), which suggests a calculation of the frequency of how many times a specific word appears in the email, expressed as a percentage of total words in the email multiplied by 100.
- Some variables are named char_freq_(number), which suggests a count of the frequency of the specific ensuing character, expressed as a percentage of total characters in the email multiplied by 100. Note, these characters are not valid column names in R, but you can view them in the raw .csv file
- Some variables are named capital_run_length_(number) which suggests some information about the average (or maximum length of, or total) consecutive capital letters in the email.
- spam: This is the response variable, 0 = not spam, 1 = spam.

2.0.2 Missing Values

Unfortunately, the capital run length average variable is corrupted and as a result, contains a fair number of missing values. These show up as NaN (the default way of representing missing values in Python.)

2.1 Part a

Use k-nearest neighbors regression with k=15 to impute the missing values in the $capital_run_length_average$ column using the other predictors after standardizing (i.e. rescaling) them. You may use a function such as KNeighborsRegressor from the package sklearn.neighbors that performs k-nearest neighbors regression. There is no penalty for using a built-in function

When you are done with this part, you should have no more NaN's in the capital run length average column in either the training or the test set. To keep the training and test sets separate, you will need to build two models for imputing; one that is trained on, and imputes for, the training set, and another that is trained on, and imputes for, the test set. Make sure you show all of your work. (You may find the function np.isnan() useful for this problem.)

```
In [177]:  # my notes
# regression model: codomain of model is a continuous space, e.g. R
# classification model: codomain of model is a discrete space, e.g. {0,1}
In [178]: train = pd.read_csv("spam_train.csv")
                test = pd.read_csv("spam_test.csv")
```

2.1.1 Scaling

```
In [179]: sc = StandardScaler()
In [180]: def scaleAllColumnsExceptCRLA(df, sc, fit_type):
    # drop the column that remains unscaled
                    CRLA_column = df[["capital_run_length_average"]]
                         = df.drop(columns="capital_run_length_average")
                    if fit_type
                          fit_type == "fit_transform":
scaled_columns = sc.fit_transform(df)
                    else:
                     scaled_columns = sc.transform(df)
# create new dataframe with scaled column
                    df = pd.DataFrame(scaled_columns, index=df.index, columns=df.columns)
                    # concatenate and return
return pd.concat([CRLA_column, df], axis=1)
In [181]: train_scaled = scaleAllColumnsExceptCRLA(train, sc, "fit_transform")
test_scaled = scaleAllColumnsExceptCRLA(test, sc, "fit_transform")
```

2.1.2 Fit KNR to non-NaN frames

1 of 8

```
Contents 2 4
 1 Assignment 2
▼ 2 Problem 1: Spam, wonderful spam!
      2.0.1 A Note about Features
      2.0.2 Missing Values
  ▼ 2.1 Part a
      2.1.1 Scaling
      2.1.2 Fit KNR to non-NaN frames
    2.2 Part b
  ▼ 2.3 Part c
      2.3.1 KNN - all predictors except ca
      2.3.2 KNN - all predictors
      2.3.3 Logistic Regression - all predic
    ▼ 2.3.4 Logistic Regression - all predic
        2.3.4.1 Significant predictors
      2.3.5 Export to CSV
▼ 3 Problem 2: Gradient Descent
    3.1 Part a
    3.2 Part b
    3.3 Part c
    3.4 Part d
▼ 4 Problem 3: Cross-Validation
    4.1 Part a
    4.2 Part b
    4.3 Part c
    4.4 Part d
    4.5 Part e
    4.6 Part f
```

```
In [182]: from sklearn.neighbors import KNeighborsRegressor as knr
In [183]: def fitKNRToData(notNaNData):
    regressor = knr(n_neighbors=15)
                    X = notNaNData.drop(columns="capital_run_length_average")
y = notNaNData["capital_run_length_average"]
regressor.fit(X, y)
                    return regressor
In [184]: def createPredictionsFromRegressor(regressor, NaNData):
    predictions = regressor.predict(NaNData.drop(columns="capital_run_length_average"))
                     return predictions
In [185]: def combinePredictionsWithNaNData(predictions, NaNData):
                    NaNData = NaNData.drop(columns="capital_run_length_average")
NaNData["capital_run_length_average"] = predictions
                     return NaNData
In [186]: def imputePipeline(df):
                    imputerspeline(ar):
# split into NaN and non-NaN rows
not_nan_data = df.dropna()
nan_data = df[np.isnan(df.capital_run_length_average)]
                      # fit KNR to non-NaN row
                    regressor = fitKNRToData(not_nan_data)
                    # make predictions
predictions = createPredictionsFromRegressor(regressor, nan_data)
                    # add in to the previously NaN data
imputed = combinePredictionsWithNaNData(predictions, nan_data)
                     # sort imputed and non_nan
imputed = imputed.reindex(sorted(imputed.columns), axis=1)
                     not_nan_data = not_nan_data.reindex(sorted(not_nan_data.columns), axis=1)
                    rejoin = pd.concat([not_nan_data, imputed])
return rejoin
In [187]: train_scaled_imputed = imputePipeline(train_scaled)
test_scaled_imputed = imputePipeline(test_scaled)
```

2.2 Part b

Write a function named knnclass() that performs k-nearest neighbors classification, without resorting to a package. Essentially, we are asking you to recreate the sklearn.neighbors.KNeighborsClassifier function; though, we do not expect you to implement a fancy nearest neighbor search algorithm like what KNeighborsClassifier uses, just the naive search will suffice. Additionally, this function will be more sophisticated in the following way:

- The function should automatically do a split of the training data into a sub-training set (80%) and a validation set (20%) for selecting the optimal k.(More sophisticated cross-validation is not necessary.)
- The function should standardize each column: for a particular variable, say x_1 , compute the mean and standard deviation of x_1 using the training set only, say \bar{x}_1 and s_1 ; then for each observed x_1 in the training set and test set, subtract \bar{x}_1 , then divide by s_1 .

Note: You can assume that all columns will be numeric and that Euclidean distance is the distance measure.

The function skeleton is provided below.

2 of 8 9/26/19, 11:52 PM

```
Contents 2 4
  1 Assignment 2
▼ 2 Problem 1: Spam, wonderful spam!
      2.0.1 A Note about Features
      2.0.2 Missing Values
  ▼ 2.1 Part a
      2.1.1 Scaling
      2.1.2 Fit KNR to non-NaN frames
    2.2 Part b
  ▼ 2.3 Part c
      2.3.1 KNN - all predictors except ca
      2.3.2 KNN - all predictors
      2.3.3 Logistic Regression - all predic
     ▼ 2.3.4 Logistic Regression - all predic
        2.3.4.1 Significant predictors
      2.3.5 Export to CSV
▼ 3 Problem 2: Gradient Descent
    3.1 Part a
    3.2 Part b
    3.3 Part c
     3.4 Part d
 ▼ 4 Problem 3: Cross-Validation
    4.1 Part a
    4.2 Part b
    4.3 Part c
     4.4 Part d
    4.5 Part e
    4.6 Part f
```

```
In [211]: class KnnClass:
                       def __init__(self, xtrain, ytrain):
                              # training data
                              self.x train = xtrain
                             self.y_train = ytrain
                # do a further split of the training data
# so x train_train, y_train_train to train model
# x_train_test, y_train_test to validate model
self.x_train_test, y_train_test to validate model
self.x_train_train, self.x_train_test, self.y_train_train, self.y_train_test = train_test_split(self.x_train, self.y_train, test_size=0.2, random_state=42)
                             \# after fit() is called, this is updated self.optimal_k = -1
                       def euclideanDistance(self, arr1, arr2):
                              return np.linalg.norm(arr1 - arr2)
                       def runTests(self):
                             print(self.euclideanDistance(self.x_train.iloc[0], self.x_train.iloc[43]))
                          dataPoint: Series or Numpy Array
                       def getMeighbors(self, dataSet, dataPoint, kNeighbors):
    distances = []
    # iterate through the data set, calculating distances to all neighbors
                             # iterate through the data set, calculating distances
for index in range(len(dataSet)):
    currRow = dataSet.iloc[index]
    dist = self.euclideanDistance(currRow, dataPoint)
    distances.append({
        "dist": dist,
        "index": index
                             \label{limits}  \mbox{distances.sort(key=lambda x: x["dist"]) \# sort \ by \ distance return \ distances[:kNeighbors] }
                       def classifyUsingNeighbors(self, labelsData, neighbors):
    # keep counters to record votes
    spam = 0
                                ot_spam = 0
                                 iterate through the neigbors
                             for n in neighbors:
    index = n["index"]
    # check the index against the labels
                                    if labelsData.iloc[index] > 1.0:
    spam += 1
                                    else:
                             not_spam += 1
return 1 if spam > not_spam else 0
                       def getAccuracy(self, trueLabels, predictedLabels):
    accurate = 0
    total = len(trueLabels)
                              for i in range(len(trueLabels)):
                                    if trueLabel[i] and predictedLabel[i] >= 1:
    accurate += 1
                                    elif trueLabel[i] and predictedLabel[i] < 1:</pre>
                             accurate +=1
return accurate / total
                              # find the best k from 1:15 using the sub training data
                              accuracies = []
for kNeighbors in range(1, 15):
                                    # classify
                                    for index in range(len(self.x_train_test)):
    dataPoint = self.x_train_test.iloc[index]
    neighbors = self.getNeighbors(self.x_train_train, dataPoint, kNeighbors)
                                           label = self.classifyUsingNeighbors(self.y_train_train, neighbors)
                                      spamLabels.append(label)
# record accuracy for this K
                                    accuracies.append({
    "k" : kNeighbors,
                                           "accuracy" : getAccuracy(self.y train test, spamLabels)
                              # update k with best accuracy
                             accuracies.sort(key=lambda a: a["accuracy"], reverse=True)
self.optimal_k = accuracies["k"]
                      def predict(self, x_test, k=-1):
    k = self.optimal_k if k < 0 else k
    spamLabels = []
    for index in range(len(x_test)):</pre>
                                    adataPoint = x_test.iloc[index]
neighbors = self.getNeighbors(self.x_train, dataPoint, k)
label = self.classifyUsingNeighbors(self.y_train, neighbors)
                             spamLabels.append(label)
return spamLabels
In [198]: KC = KnnClass(x_train_knn, y_train_knn)
# KC.fit()
                # KC.predict(x_test_knn)
```

3 of 8 9/26/19, 11:52 PM

Contents 2 4

- 1 Assignment 2
- ▼ 2 Problem 1: Spam, wonderful spam! 2.0.1 A Note about Features 2.0.2 Missing Values
- ▼ 2.1 Part a
- 2.1.1 Scaling
- 2.1.2 Fit KNR to non-NaN frames 2.2 Part b
- ▼ 2.3 Part c
- 2.3.1 KNN all predictors except ca 2.3.2 KNN - all predictors 2.3.3 Logistic Regression - all predic
- ▼ 2.3.4 Logistic Regression all predic 2.3.4.1 Significant predictors 2.3.5 Export to CSV
- ▼ 3 Problem 2: Gradient Descent
- 3.1 Part a 3.2 Part b
- 3.3 Part c
- 3.4 Part d
- ▼ 4 Problem 3: Cross-Validation
 - 4.1 Part a
 - 4.2 Part b
 - 4.3 Part c
 - 4.6 Part f
- 4.4 Part d 4.5 Part e

- Timer unit: 1e-06 s Total time: 74.5331 s File: <ipython-input-197-771ec6685ebf> Function: fit at line 74 Line # Time Per Hit % Time Line Contents 74
 - def fit(self):
 # find the best k using the split data 75 accuracies = [] for kNeighbors in range(1, 51): 76 77 78 2.0 2.0 0.0 spamLabels = [] for index in range(len(self.x_train_test) - 594):
 dataPoint = self.x_train_test.iloc[index]
 neighbors = self.getNeighbors(self.x_train_trai 42.0 12210.0 0.8 79 51 0.0 50 80 0.0 81 50 74505510.0 1490110.2 100.0 n, dataPoint, 15) 15319.0 0.0 306.4 label = self.classifyUsingNeighbors(self.y train 82 50 train, neighbors) 83 50 spamLabels.append(label) 84 1 0.0 0.0 0.0 return spamLabels
- Naive search for neighbors is too slow on a training set of size 3220.
- . So for Part C, I will be using a built-in package which utilize trees and cache for fast search.
- My Class works, but it takes 12 min for a single K, so utilizing it for all K is not feasible.
- You can take a look at the profiler, and you will see that getting the neighbors takes close to 100% of the time.
- It takes 1 minute on my laptop to classify 50 points. By this estimate, fitting the model for single K will take around 15 minutes.

I tried till k = 5, and my optimal k was k == 5

In [199]: %lprun -f KC.fit KC.fit()

2.3 Part c

In this part, you will need to use a k-NN classifier to fit models on the actual dataset. If you weren't able to successfully write a k-NN classifier in Part b, you're permitted to use a built-in package for it. If you take this route, you may need to write some code to standardize the variables and select k, which knnclass() from part b already does

Now fit 4 models and produce 4 sets of predictions of spam on the test set:

- 1. knnclass() using all predictors except for capital_run_length_average (say, if we were distrustful of our imputation approach). Call these predictions knn_pred1
- 2. knnclass() using all predictors including capital_run_length_average with the imputed values. Call these predictions knn_pred2
- $3. \ logistic \ regression \ using \ all \ predictors \ except \ for \ capital_run_length_average \ . \ Call \ these \ predictions \ logm_pred1 \ .$
- $4. \ logistic regression using all predictors including \ {\tt capital_run_length_average} \ with the imputed values. \ Call these predictions \ {\tt logm_pred2} \ .$

In 3-4 sentences, provide a quick summary of your second logistic regression model (model 4). Which predictors appeared to be most significant? Are there any surprises in the predictors that ended up being significant or not significant?

Submit a .csv file called assn2 NETID results.csv that contains 5 columns:

- capital_run_length_average : the predictor in your test set that now contains the imputed values (so that we can check your work on imputation).
- knn_pred1
- knn_pred2
- logm_pred1
- logm_pred2

Make sure that row 1 here corresponds to row 1 of the test set, row 2 corresponds to row 2 of the test set, and so on.

2.3.1 KNN - all predictors except capital_run_length_average

```
In [257]: knnClassifier = KNeighborsClassifier(n neighbors=5)
          knnClassifier.fit(x_train_knn.drop(columns=["capital_run_length_average"]).values, [1 if i > 1 else 0 for i in
          y train knn.values])
          knn_pred1 = knnClassifier.predict(x_test_knn.drop(columns=["capital_run_length_average"]).values)
```

2.3.2 KNN - all predictors

```
In [258]: knnclassifier2 = KNeighborsClassifier(n_neighbors=5)
knnclassifier2.fit(x_train_knn.values, [1 if i > 1 else 0 for i in y_train_knn.values])
              knn_pred2 = knnClassifier2.predict(x_test_knn.values)
```

2.3.3 Logistic Regression - all predictors except capital_run_length_average

```
In [259]: lrClassifier = LogisticRegression()
          lrClassifier.fit(x_train_knn.drop(columns=["capital_run_length_average"]).values, [1 if i > 1 else 0 for i in y
          logm_pred1 = lrClassifier.predict(x_test_knn.drop(columns=["capital_run_length_average"]).values)
```

/Users/sarimabbas/Developer/dataScience/sds355_container/sds355_env/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning. FutureWarning)

2.3.4 Logistic Regression - all predictors

9/26/19, 11:52 PM 4 of 8

0.195161

0.351862

0.099003

0.173792

-0.328363

1.958048

-0.499310

0.598068

-0.164188

-0.164188

-0.164188

-0.164188 -0.

0.115010

-0.164188

-0 164188

0.323371 -0.164188 -0.

-0.

-0.

```
In [269]: lrClassifier2 = LogisticRegression()
Contents 2 4
                                                                           lrModel2 = lrClassifier2.fit(x_train_knn.values, [1 if i > 1 else 0 for i in y_train_knn.values])
logm_pred2 = lrClassifier2.predict(x_test_knn.values)
   1 Assignment 2
▼ 2 Problem 1: Spam, wonderful spam!
       2.0.1 A Note about Features
                                                                           /Users/sarimabbas/Developer/dataScience/sds355_container/sds355_env/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to
       2.0.2 Missing Values
   ▼ 2.1 Part a
                                                                           silence this warning.
                                                                               FutureWarning)
       2.1.1 Scaling
       2.1.2 Fit KNR to non-NaN frames
   ▼ 2.3 Part c
                                                          2.3.4.1 Significant predictors
       2.3.1 KNN - all predictors except ca
        2.3.2 KNN - all predictors
                                                           In [281]: x_train_knn
        2.3.3 Logistic Regression - all predic
     ▼ 2.3.4 Logistic Regression - all predic
                                                           Out[281]:
          2.3.4.1 Significant predictors
                                                                                     capital_run_length_average capital_run_length_longest capital_run_length_total char_freq_! char_freq_# char_freq_$ char_freq_( char_freq_;
        2.3.5 Export to CSV
                                                                                0
                                                                                                                                        -0.027690
                                                                                                                                                                     1.247501
                                                                                                                                                                                  -0.303582
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 -0.266994
▼ 3 Problem 2: Gradient Descent
                                                                                                         1.687000
                                                                                                                                        -0.199141
                                                                                                                                                                    -0.404020 -0.034317
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 -0.337192
     3.1 Part a
     3.2 Part b
                                                                                2
                                                                                                        1.750000
                                                                                                                                        -0.213042
                                                                                                                                                                    -0.347180 -0.114308
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 -0.337192
     3.3 Part c
                                                                                                        5.038000
                                                                                3
                                                                                                                                        0.032549
                                                                                                                                                                    -0.239815
                                                                                                                                                                                -0.197678
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 1.876390
     3.4 Part d

▼ 4 Problem 3: Cross-Validation

                                                                                                                                        -0.175972
                                                                                                         1.785000
                                                                                                                                                                    -0.328233 -0.303582
                                                                                                                                                                                                  -0.126696
                                                                                                                                                                                                                 -0.337192
     4.1 Part a
     4.2 Part b
                                                                                                                                        -0.171338
                                                                                                                                                                    -0.219289 -0.303582 -0.126696
                                                                                                                                                                                                                 -0.337192
     4.3 Part c
                                                                             3200
                                                                                                        2.098133
     4.4 Part d
                                                                                                                                                                                -0.285556
                                                                             3207
                                                                                                                                        -0.111099
                                                                                                                                                                    0.088594
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 -0.037680
                                                                                                        2.349533
     4.5 Part e
                                                                                                                                        -0.083296
                                                                                                                                                                                                                 0.252472
     4.6 Part f
                                                                                                        2.968800
                                                                                                                                        0.083520
                                                                                                                                                                     0.447003
                                                                                                                                                                                   0.064827
                                                                                                                                                                                                  -0.126696
                                                                                                                                                                                                                 -0.337192
                                                                             3212
                                                                             3217
                                                                                                        1 555933
                                                                                                                                        -0.231577
                                                                                                                                                                    -0.437176 -0.303582
                                                                                                                                                                                                 -0.126696
                                                                                                                                                                                                                 -0.337192
                                                                           3220 rows × 57 columns
                                                           In [300]:
                                                                           coeffs = []
for index, cf in enumerate(lrModel2.coef_[0]):
                                                                                  coeffs.append({
    "index" : index,
    "cf" : cf
                                                                            coeffs.sort(key=lambda x : x["cf"], reverse=True)
                                                           In [301]: coeffs
                                                          'index': 49, 'cf': 0.7802056135918461},
'index': 28, 'cf': 0.7452619483387634},
'index': 11, 'cf': 0.7424063561987302},
                                                                                index': 4, 'cf': 0.6443895856549169},
'index': 21, 'cf': 0.5793399173060575},
'index': 38, 'cf': 0.43386436294173547},
                                                                                                    'cf': 0.42524857740901895},
'cf': 0.3551553040203605},
                                                                                 index': 19,
                                                                                'index': 41,
                                                                                Index: 41, c1: 0.33313340429335827),
'index': 56, 'cf': 0.3333377244536883),
'index': 27, 'cf': 0.2986137669190685),
'index': 27, 'cf': 0.2771960391204137),
'index': 17, 'cf': 0.27671800946342234),
'index': 32, 'cf': 0.24964294203130896),
                                                                                                  'cf': 0.23769801745698302},
'cf': 0.18515181862471466},
'cf': 0.15383402259570236},
                                                                                'index': 3, 'index': 42,
                                                                                 index': 13.
                                                                                index': 55,
                                                                                                    'cf': 0.13989332669603638}, 'cf': 0.12274363237634585},
                                                                                'index': 35,
                                                                                index': 26,
index': 18,
                                                                                                    'cf': 0.1066451288376731},
'cf': 0.10445741975210442},
'cf': 0.07452521026414054},
                                                                                index': 39,
                                                                                'index': 0, 'cf': 0.07375305989089152},
'index': 50, 'cf': 0.04517879625872991,
'index': 6, 'cf': 0.01837790684119818},
                                                                                                    'cf': 0.000325806581166473},
'cf': -0.04718301433236289},
                                                                                index': 48,
                                                                                'index': 44,
                                                                                                    'cf': -0.05323241519665686},
'cf': -0.0603959177836788},
                                                                                'index': 36,
                                                                                                    'cf': -0.0749952542140989},
                                                                                index': 54,
                                                                                                    'cf': -0.08898849450571826},
'cf': -0.10800028917584115},
'cf': -0.11329307077837648},
                                                                                index': 12,
                                                                                index': 43.
                                                                                                    'cf': -0.11834382781751411},
                                                                                                  'cf': -0.15032089437332502},
'cf': -0.15104702422603858},
'cf': -0.21805066262971576},
                                                                                index': 40,
                                                                                'index': 8, 'index': 15,
                                                                                index': 16, 'cf': -0.2321104941102357},
                                                                                                   'cf': -0.26341591150844496}
'cf': -0.2949123739999225},
                                                                                'index': 23, 'cf': -0.4215542945405113}, 'index': 45, 'cf': -0.4634880016191258},
                                                                                index': 53,
                                                                                                    'cf': -0.5863738148024542},
                                                                                                    'cf': -0.6203610088575519},
'cf': -0.7956938888488201},
                                                                                 index': 14,
                                                                                index': 47.
                                                                                                    'cf': -0.828142034847825).
                                                                                'index': 20,
'index': 31,
                                                                                                    'cf': -0.8946151151870269},
'cf': -0.8996963233425389},
                                                                                index': 33,
                                                                                                    'cf': -0.9572149549841147}.
                                                                                                    'cf': -1.1292032323796697},
                                                                               'index': 25, 'cf': -1.1759474676644763}, 
'index': 37, 'cf': -1.3212986174453196}, 
'index': 30, 'cf': -2.2036224797398627},
                                                                              {'index': 29, 'cf': -3.48890040621353}]
```

9/26/19, 11:52 PM 5 of 8

Contents 2 4

```
1 Assignment 2
```

- ▼ 2 Problem 1: Spam, wonderful spam! 2.0.1 A Note about Features 2.0.2 Missing Values ▼ 2.1 Part a 2.1.1 Scaling 2.1.2 Fit KNR to non-NaN frames 2.2 Part b ▼ 2.3 Part c
 - 2.3.1 KNN all predictors except ca 2.3.2 KNN - all predictors 2.3.3 Logistic Regression - all predic ▼ 2.3.4 Logistic Regression - all predic
 - 2.3.4.1 Significant predictors 2.3.5 Export to CSV
- ▼ 3 Problem 2: Gradient Descent
- 3.1 Part a 3.2 Part b 3.3 Part c
- 3.4 Part d
- ▼ 4 Problem 3: Cross-Validation
 - 4.1 Part a 4.2 Part b
 - 4.3 Part c 4.4 Part d
 - 4.5 Part e 4.6 Part f

Based on these standardized coefficients, significant predictors include:

- · capital run length total
- charfreq\$
- word_freq_000

2.3.5 Export to CSV

```
In [261]: exportDict = {
    "capital_run_length_average" : x_test_knn["capital_run_length_average"].values,
    "knn_pred1" : knn_pred2,
    "logm_pred1" : logm_pred2,
    "logm_pred2" : logm_pred1,
    "logm_pred2" : logm_pred2
}
In [263]: exportFrame = pd.DataFrame(exportDict)
In [266]: # exportFrame.to_csv("./file.csv")
```

3 Problem 2: Gradient Descent

Consider the scenario of univariate logistic regression where we are trying to predict Y, which can take the value 0 or 1, from the variable X, which can take the value of any real number. Recall from lecture that we need to predict parameters β_0 and β_1 by minimizing the penalized loss function:

$$L(\beta_0,\beta_1) = \sum_{i=1}^n \left[log \left(1 + e^{\beta_0 + X_i \beta_1} \right) - Y_i \left(\beta_0 + X_i \beta_1 \right) \right] + \lambda \left(\beta_0^2 + \beta_1^2 \right).$$

Run the next cell to simulate data from the true values of β_0 = 2.5 and β_1 = 3.0.

```
In [212]: n = 10000
            x1 = np.random.uniform(-5, 5, size=n)
           beta0 = 2.5
beta1 = -3.0
p = np.exp(beta0 + x1*beta1)/(1 + np.exp(beta0 + beta1*x1))
           y = np.random.binomial(1, p, size=n)
In [215]: display(x1)
           display(y)
           array([ 3.6447943 , -1.77319003, 1.70788791, ..., 4.44073042, -3.98045186, 0.82309827])
           array([2.17219945e-04, 9.99598367e-01, 6.76309040e-02, ...,
                   1.99514009e-05, 9.99999465e-01, 5.07675693e-01])
           array([0, 1, 0, ..., 0, 1, 1])
```

3.1 Part a

For given values of β_0 and β_1 the vector $\left(\frac{\partial}{\partial \beta_0} L(\beta_0, \beta_1), \frac{\partial}{\partial \beta_1} L(\beta_0, \beta_1)\right)^T$ is called the gradient of $L(\beta_0, \beta_1)$ and is denoted $\nabla L(\beta_0, \beta_1)$.

Calculate the derivative of $L(\beta_0, \beta_1)$ with respect to β_0 , treating β_1 as a constant. (i.e. calculate $\frac{\partial}{\partial \beta_0} L(\beta_0, \beta_1)$).

Now calculate the derivative of $L(\beta_0, \beta_1)$ with respect to β_1 , treating β_0 as a constant. (i.e. calculate $\frac{\partial}{\partial \beta_1} L(\beta_0, \beta_1)$).

Be sure to show your work by either typing it in here using LaTeX, or by taking a picture of your handwritten solutions and displaying them here in the notebook. (If you choose the latter of these two options, be sure that the display is large enough and legible. You may find the example shown in the Introduction to Python.ipynb notebook for the Yale image useful.)

```
In [ ]: ## please put your answer here ##
```

3.2 Part b

Complete the function in the following cell called update() which takes values for β_0 and β_1 as well as a step-size η and should return updated values for β_0 and β_1 from one step of gradient descent (using all the data and your answer to Part a). You may use the value 0.01 for λ .

```
In [7]: def update(b0, b1, eta):
```

Now complete the function in the next cell called *loss()* which takes values for β_0 and β_1 and should return the value of the loss function evaluated at those two parameter values.

```
In [8]: def loss(b0, b1):
```

3.3 Part c

Contents ₽ ♥

```
1 Assignment 2
```

- ▼ 2 Problem 1: Spam, wonderful spam! 2.0.1 A Note about Features 2.0.2 Missing Values
- ▼ 2.1 Part a
- 2.1.1 Scaling
- 2.1.2 Fit KNR to non-NaN frames 2.2 Part b
- ▼ 2.3 Part c
- 2.3.1 KNN all predictors except ca 2.3.2 KNN - all predictors
- 2.3.3 Logistic Regression all predic ▼ 2.3.4 Logistic Regression - all predic
- 2.3.4.1 Significant predictors 2.3.5 Export to CSV
- ▼ 3 Problem 2: Gradient Descent
- 3.1 Part a 3.2 Part b
- 3.3 Part c 3.4 Part d
- ▼ 4 Problem 3: Cross-Validation
 - 4.1 Part a
 - 4.2 Part b 4.3 Part c
 - 4.4 Part d
 - 4.5 Part e
 - 4.6 Part f

The following cell uses the two functions from Part b to implement gradient descent for this problem, keeping track of the values for β_0 , β_1 , and $L(\beta_0,\beta_1)$ at each iteration. In the cell below the code, answer each of the questions included as comments next to the code. Also, create individual plots of β_0 , β_1 , and $L(\beta_0,\beta_1)$ vs. iteration number. Do these three quantities behave as expected for gradient descent?

```
In [28]: step = 0.01
    beta0_hat = 0
    beta1_hat = 0
    l = loss(beta0_hat, beta1_hat)
              beta0_all = [beta0_hat]
beta1_all = [beta1_hat]
loss_all = [1]
               while i < 400 and step > 3e-8: #1. What is the reasoning behind these two stopping criteria?
b = update(beta0_hat, beta1_hat, step) #2. What is being calculated here?
                     1_new = loss(b[0], b[1])
if l_new < 1:
    beta0_hat = b[0]</pre>
                                                                                    #3. What is being calculated here?
#4. What happens if the statement being tested here is True?
                           beta1_hat = b[1]
                           1 = 1 \text{ new}
                     else:
                            step = step*0.9
                                                                                    #5. What happens if the statement tested above is False? What is th
               e reasoning
                                                                                     # behind this?
                     beta0_all.append(beta0_hat)
beta1_all.append(beta1_hat)
                     loss_all.append(1)
```

```
1.
```

2.

4

5.

In []: ## Create your plots here ##

3.4 Part d

Is your gradient descent algorithm from Part b robust against initial estimates of β_0 and β_1 ? To help answer this question, take the code above that implements gradient descent and put it in a function that takes initial estimates of β_0 and β_1 as arguments, and returns the optimized values from gradient descent. Run this function using each of the following pairs of (β_0, β_1) as initial estimates: (15, 3), (-30, 5), and (-8, -8). Are your final estimates approximately the same each time?

```
In [ ]: ## please put your answer here ##
```

4 Problem 3: Cross-Validation

4.1 Part a

Generate a simulated data set with the following cell:

```
In [205]: np.random.seed(1)
```

In this data set, what is the value of n (the number of data points) and what is the value of p (the true number of model parameters)? Write out the model used to generate the data in equation form.

```
N = 100
p = 2
```

 $y = x - 2x^2 + \epsilon$

4.2 Part b

Create a scatterplot of X against Y. Comment on what you find

7 of 8 9/26/19, 11:52 PM

Contents 2 ♥

1 Assignment 2

▼ 2 Problem 1: Spam, wonderful spam! 2.0.1 A Note about Features 2.0.2 Missing Values

▼ 2.1 Part a

2.1.1 Scaling

2.1.2 Fit KNR to non-NaN frames 2.2 Part b

▼ 2.3 Part c

2.3.1 KNN - all predictors except ca

2.3.2 KNN - all predictors
2.3.3 Logistic Regression - all predic

▼ 2.3.4 Logistic Regression - all predic 2.3.4.1 Significant predictors

2.3.5 Export to CSV ▼ 3 Problem 2: Gradient Descent

3.1 Part a 3.2 Part b

3.2 Part b 3.3 Part c

3.4 Part d

▼ 4 Problem 3: Cross-Validation

4.1 Part a

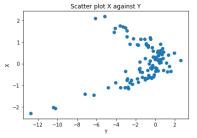
4.2 Part b

4.3 Part c 4.4 Part d

4.5 Part e

4.6 Part f

In [13]:
 plt.scatter(y, x) # questions wants X against Y i.e. X on the y-axis
 plt.title('Scatter plot X against Y')
 plt.xlabel('Y')
 plt.ylabel('X')
 plt.show()



The graph looks quadratic, with Y maximum near $X=\mathbf{0}$

4.3 Part c

Set a random seed, and then compute the Leave-One-Out Cross-Validation (LOOCV) errors that result from fitting the following four models using least squares:

i.
$$Y = \beta_0 + \beta_1 X + \epsilon$$

ii. $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \epsilon$

iii. $Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$

iv.
$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \beta_4 X^4 + \epsilon$$

Note: for linear regression, the LOOCV error can be computed via the following short-cut formula:

LOOCV Error =
$$\frac{1}{n} \sum_{i=1}^{n} \left(\frac{Y_i - Y_i}{1 - H_{ii}} \right)^{i}$$

where H_{ii} is the i^{th} diagonal entry of the projection matrix $H = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T$, and \mathbf{X} is a matrix of predictors (the design matrix). This formula is an alternative to actually carrying out the n = 100 regressions you would otherwise need for LOOCV. An example of how to calculate the projection matrix H is provided below for the case of n = 5 and the model $Y = \beta_0 + \beta_1 X + \beta_2 Y^2 + \epsilon$. To get the diagonal elements of H you may find the function np.diag() useful.

```
In [16]: example_x = np.array([-3, -4, -5, -6, -7])  #generate the x variable
    design_x = np.vander(example_x, 3)  #calculate the design matrix for the polynomial model with 3 fit pa
    rameters
    H = np.dot(design_x, np.dot(np.linalg.inv(np.dot(design_x.T, design_x)), design_x.T))  #calculate H
In []: ## please put your answer here ##
```

4.4 Part d

Repeat Part c using another random seed to generate data, and report your results. Are your results the same as what you got in Part c? Why?

In [18]: ## please put your answer here ##

4.5 Part e

Which of the models in Part c had the smallest LOOCV error? Is this what you expected? Explain your answer.

In [19]: ## please put your answer here ##

4.6 Part f

Comment on the statistical significance of the coefficient estimates that results from fitting each of the models in Part c using least squares. Do these results agree with the conclusions drawn based on the cross-validation results?

In [20]: ## please put your answer here ##

8 of 8 9/26/19, 11:52 PM