An aerial photograph of a vast, rugged mountain range covered in snow. The peaks are jagged and partially obscured by shadows, creating a high-contrast scene. The snow appears thick and uneven, with some rocky outcrops visible. The overall tone is cool and serene, with a mix of white, grey, and blue hues.

S&DS 355 / 365 / 565  
**Data Mining and Machine Learning**

# **Classification and Gradient Descent**

Tuesday, September 10

Yale

# Office Hours

365:

- Derek Feng: Tuesday 5:00pm-6:00pm, DL #225
- Brandon Chow: Mondays, 6:30pm-7:30pm, 24 Hillhouse (Classroom)
- Colleen Chan: Thursdays 5:00-6:00pm, 17 Hillhouse Rm 7 (Lower Level)
- Kasra Esfandiari: Wednesdays, 6:30pm-7:30pm, 17 Hillhouse Rm 220
- (ULA) Qinying Sun: Friday, 2:00pm-3:00pm, 17 Hillhouse Rm 110

# Office Hours

355:

- Parker Holzer: Wednesday 4:00-5:00pm (24 Hillhouse Ave, classroom 107)
- Xinyi Zhong: Thursday 5:00-6:00pm (24 Hillhouse Ave, classroom 107)
- Zifan Li: Monday 7:30 - 8:30pm (24 Hillhouse Ave, classroom 107)

ULAs:

- Max Yuhas: Wednesday 6:00-7:00pm (17 Hillhouse, Room 07)
- Chloe Zhou: Thursday 7:30-8:30pm (Bass L06-A, for September 12)
- Daniel Zhou: Monday TBD
- Adriel Sumathipala: TBD

# Outline

- Shrinkage, bias and variance
- Regularization
- Stochastic gradient descent

# What did we talk about last time?

- Logistic regression is a linear model of the log-odds
- If data are perfectly separable, log-likelihood will be unbounded

# Penalization, shrinkage, bias and variance

Estimator  $\hat{\theta}$  of a parameter  $\theta$ :

$$\text{bias}^2 \quad \left( \mathbb{E}(\hat{\theta}) - \theta \right)^2$$

$$\text{variance} \quad \mathbb{E}(\hat{\theta} - \mathbb{E}(\hat{\theta}))^2$$

# Penalization, shrinkage, bias and variance

Estimator  $\hat{\theta}$  of a parameter  $\theta$ :

$$\begin{array}{ll} \text{bias}^2 & \left( \mathbb{E}(\hat{\theta}) - \theta \right)^2 \\ \text{variance} & \mathbb{E}(\hat{\theta} - \mathbb{E}(\hat{\theta}))^2 \end{array}$$

Expected squared error decomposes as

$$\mathbb{E}(\hat{\theta} - \theta)^2 = \text{bias}^2 + \text{variance}$$

# Penalization, bias and variance

Let's see how shrinkage affects the bias and variance.

Suppose  $Y \sim N(\theta, \sigma^2)$ .

(a)  $\hat{\theta} = Y$ . Bias? Variance?



# Penalization, bias and variance

Let's see how shrinkage affects the bias and variance.

Suppose  $Y \sim N(\theta, \sigma^2)$ .

(a)  $\hat{\theta} = Y$ . Bias? Variance?

(b)  $\hat{\theta} = bY$ , for  $0 \leq b \leq 1$ . Bias? Variance?

# Exercise

Consider the simplified version of the objective function

$$F(\beta) = (Y - \beta)^2 + \lambda\beta^2$$

What is the minimizer  $\hat{\beta}$ ?

# Exercise

Consider the simplified version of the objective function

$$F(\beta) = (Y - \beta)^2 + \lambda\beta^2$$

What is the minimizer  $\hat{\beta}$ ?

$$\hat{\beta} = \left( \frac{1}{1 + \lambda} \right) Y$$

# Exercise

Now let's add a predictor variable,

$$F(\beta) = (Y - X\beta)^2 + \lambda\beta^2$$

What is the minimizer?

# Exercise

Now let's add a predictor variable,

$$F(\beta) = (Y - X\beta)^2 + \lambda\beta^2$$

What is the minimizer?

$$\hat{\beta} = \frac{XY}{X^2 + \lambda}$$

# Shrinkage

In a *shrinkage estimator*, we squash down the estimate by a scaling factor. For example,

$$\hat{\beta} \leftarrow \left( \frac{1}{1 + \lambda} \right) \hat{\beta}$$

This induces bias-variance tradeoff — the bias goes up, but the variance goes down.

# Penalization

To guard against overfitting, we can *penalize* the coefficients:

$$F(\beta) = -\log\text{-likelihood}(\beta) + \lambda \|\beta\|^2$$

- Large coefficients incur a large penalty
- The *regularization parameter*  $\lambda$  controls the tradeoff between fit to the data, and size of the coefficients
- Small  $\lambda$ : high variance, low bias
- Large  $\lambda$ : low variance, high bias
- As  $\lambda$  increases, the size of the coefficients  $|\beta_j|$  decreases.

# Political blog classification

- Political Blog Classification. A collection of 403 political blogs were collected during two months before the 2004 presidential election. The goal is to predict whether a blog is *liberal* ( $Y = 0$ ) or *conservative* ( $Y = 1$ ) given the content of the blog.







# Political blog classification

- 403 blogs
- 205 are “liberal” and 198 are “conservative”
- For each word, value of a feature is word frequency
- Lower case and remove highly frequent words, throw out those appearing fewer than 10 times.
- 23,955 features
- Links to 292 popular blogs included as binary vector

# Political blog classification results




without links




with links

# Political blog classification results



# Newton's method




- “Grand-daddy of optimization”
- Fast convergence
- Great properties
- Second order method
- Not scalable

# Gradient Descent

**Gradient descent** is a procedure for finding the arguments that minimize a particular function (called a **cost function**).

e.g. cost function could be a negative likelihood or negative log-likelihood.



# Gradient Descent


Goal: Find  $(\theta_1, \dots, \theta_p)$  that minimizes **cost function**  $L(\theta_1, \dots, \theta_p)$

Update equation:

$$\theta_j \leftarrow \theta_j - \rho \frac{\partial L}{\partial \theta_j}$$

In matrix form:

$$\theta \leftarrow \theta - \rho \nabla L(\theta)$$



$\rho$  is called the learning rate.

don't take too large steps or otherwise risk overshoot  
not too small, otherwise too slow


# Intuition for GD

355 skip

The gradient gives the direction of *steepest ascent*. Consider the Taylor Series of  $f$  about  $\theta$ ,

$$f(\theta + \delta) = f(\theta) + \nabla f(\theta) \cdot \delta + \dots,$$

for some small increment  $\delta$ . Then, ignoring higher order terms, it is clear that to maximize  $f(\theta + \delta)$  we must pick  $\delta$  in the direction of  $\nabla f(\theta)$ .



# GD for Logistic Regression

Cost function (negative log-likelihood):

$$L(\beta) = - \sum \left( y_i x_i^t \beta - \log(1 + e^{x_i^t \beta}) \right)$$

Update for logistic regression:

$$\beta_j \leftarrow \beta_j - \rho \frac{\partial L}{\partial \beta_j}.$$

$$\beta_j \leftarrow \beta_j - \rho \sum_i \left( \frac{e^{x_i^t \beta}}{1 + e^{x_i^t \beta}} - y_i \right) x_{ij}.$$



# GD for Linear Regression

Gradient descent also works for linear regression:

$$L(\beta) = \|Y - X\beta\|^2$$

$$\frac{\partial L(\beta)}{\partial \beta_j} = -2 \sum_{i=1}^n (y_i - x_i^t \beta) x_{ij}$$

GD update step:

$$\beta_j \leftarrow \beta_j + \rho \sum_{i=1}^n (y_i - x_i^t \beta) x_{ij}.$$

# GD for Linear Regression


Gradient descent also works for linear regression:

$$L(\beta) = \|Y - X\beta\|^2$$

$$\frac{\partial L(\beta)}{\partial \beta_j} = -2 \sum_{i=1}^n (y_i - x_i^t \beta) x_{ij}$$

GD update step:

$$\beta_j \leftarrow \beta_j + \rho \sum_{i=1}^n (y_i - x_i^t \beta) x_{ij}.$$



# Improving Upon Gradient Descent

Each step of (batch) gradient descent requires a calculation involving all of the data points.

# Improving Upon Gradient Descent

Each step of (batch) gradient descent requires a calculation involving all of the data points.

**Stochastic gradient descent**, in contrast, only computes based on a smaller subset of the data points (e.g. 1 observation) at each step.

## II: Stochastic gradient descent

- Suppose that we want to fit a really big model,  $n$  and  $d$  very large
- What is complexity of least squares and IRLS?

## II: Stochastic gradient descent

- Suppose that we want to fit a really big model,  $n$  and  $d$  very large
- What is complexity of least squares and IRLS?

$$O(nd^2 + d^3)$$

- Not practical for large problems
- Second order method, using covariance matrix and Hessian, solving linear system
- How can we get this to scale?

# Typical example

- We want to classify documents according to whether or not they are about corporate news.
- There are about 780,000 documents in the collection
- 60,000,000 words
- Document represented as sparse tf-idf vector

1 | 5 : 1.1789641e-01    39 : 6.0373064e-02    45 : 1.3163488e-01

- The text takes about 1.1 Gbytes
- How can we efficiently train a classifier?

# Online learning

We will introduce a method that

- Reads in the documents one at a time
- Updates the model for each document
- Updates are linear in the size of the document
- Uses little memory, never reads in the entire corpus
- Only stores a dictionary of feature weights



# Stochastic gradient descent

We initialize all weights to zero:  $\beta_j = 0, j = 1, \dots, d$ .

We read through the data one record at a time, and update the model.

- 1 Read data item  $x$
- 2 Make a prediction  $\hat{y}(x) = \sum_{j=1}^d \beta_j x_j$
- 3 Observe the true response/label  $y$
- 4 Update the weights  $\beta$  so  $\hat{y}$  is closer to  $y$

# Stochastic gradient descent

To begin, suppose we are just doing linear regression. We initialize all weights to zero:  $\beta_j = 0, j = 1, \dots, d$ .

We read through the data one record at a time, and update the model.

- 1 Read data item  $x$
- 2 Make a prediction  $\hat{y}(x) = \sum_{j=1}^d \beta_j x_j$
- 3 Observe the true response/label  $y$
- 4 Update the weights  $\beta$  so  $\hat{y}$  is closer to  $y$

$$\beta_j \longleftarrow \beta_j + \eta(y - \hat{y})x_j$$

# Stochastic gradient descent

Think about how to apply this to the data we discussed earlier


- We want to classify documents according to whether or not they are about corporate news.
- There are about 780,000 documents in the collection
- 60,000,000 words
- Document represented as sparse tf-idf vector

1 | 5 : 1.1789641e-01    39 : 6.0373064e-02    45 : 1.3163488e-01

- The text takes about 1.1 Gbytes

Key is that each data item is represented sparsely

# SGD for general loss



- Squared error  $L(y, \beta^T x) = (y - x^T \beta)^2$
- Logistic:  $L(y, \beta^T x) = -yx^T \beta + \log(1 + \exp(x^T \beta))$
- Hinge:  $L(y, \beta^T x) = (1 - yx^T \beta)_+$

# SGD for general loss

SGD update:

$$\boldsymbol{\beta} \longleftarrow \boldsymbol{\beta} - \eta \nabla L(y, \boldsymbol{\beta}^T \mathbf{x})$$

$$\beta_j \longleftarrow \beta_j - \eta \frac{\partial L(y, \boldsymbol{\beta}^T \mathbf{x})}{\partial \beta_j}$$

- $\eta$  is the *learning rate* or “step size”
- Needs to be chosen carefully, getting smaller over time

# SGD for general loss

SGD update:

$$\beta \longleftarrow \beta - \eta \nabla L(y, \beta^T x)$$

$$\beta_j \longleftarrow \beta_j - \eta \frac{\partial L(y, \beta^T x)}{\partial \beta_j}$$

Some intuition for what this is doing, and why the step size needs to decrease

$$L(\beta + \epsilon v) \approx L(\beta) + \epsilon v^T \nabla L(\beta)$$

$$L(\beta - \eta \nabla L(\beta)) \approx L(\beta) - \eta \|\nabla L(\beta)\|^2$$

This is why SGD is going downhill — if  $\eta$  is small enough

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - \pi)x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - \pi)x_j^2$$

$$\pi = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

Case checking:

- Suppose  $y = 1$  and probability  $\pi$  is high?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - \pi)x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - \pi)x_j^2$$

$$\pi = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

Case checking:

- Suppose  $y = 1$  and probability  $\pi$  is high? *small change*
- Suppose  $y = 1$  and probability  $\pi$  is small?



# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - \pi)x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - \pi)x_j^2$$

$$\pi = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

Case checking:

- Suppose  $y = 1$  and probability  $\pi$  is high? *small change*
- Suppose  $y = 1$  and probability  $\pi$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $\pi$  is small?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - \pi)x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - \pi)x_j^2$$

$$\pi = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

Case checking:

- Suppose  $y = 1$  and probability  $\pi$  is high? *small change*
- Suppose  $y = 1$  and probability  $\pi$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $\pi$  is small? *small change*
- Suppose  $y = 0$  and probability  $\pi$  is big?

# SGD for logistic regression

SGD Update:

$$\beta_j \longleftarrow \beta_j + \eta(y - \pi)x_j$$

$$\beta_j x_j \longleftarrow \beta_j x_j + \eta(y - \pi)x_j^2$$

$$\pi = \frac{1}{1 + \exp(-\beta^T \mathbf{x})}$$

Case checking:

- Suppose  $y = 1$  and probability  $\pi$  is high? *small change*
- Suppose  $y = 1$  and probability  $\pi$  is small? *big change*  $\uparrow$
- Suppose  $y = 0$  and probability  $\pi$  is small? *small change*
- Suppose  $y = 0$  and probability  $\pi$  is big? *big change*  $\downarrow$

# SGD: choice of learning rate

A conservative choice of learning rate is

$$\eta_t = \frac{1}{t}$$

A more aggressive choice is

$$\eta_t = \frac{1}{\sqrt{t}}$$

# SGD: choice of learning rate

Learning rate should scale as

$$\eta_t = \frac{1}{\sqrt{t}}$$

Problem: Some of the updates may be on different scales.

# SGD: choice of learning rate

Learning rate should scale as

$$\eta_t = \frac{1}{\sqrt{t}}$$

Problem: Some of the updates may be on different scales.

Solution: Let  $g_{tj} = \frac{\partial L(y_t, \beta^T x_t)}{\partial \beta_j}$

Scale gradients to get update rule

$$\beta_j \leftarrow \beta_j - \eta \frac{g_{tj}}{\sqrt{\sum_{s=1}^t g_{sj}^2}}$$

# SGD: scaling issues

For a linear model, the SGD update is

$$\beta_j \longleftarrow \beta_j - C_t x_j$$

If  $x_j$  increases by a factor of two, the weight  $\beta_j$  should decrease by a factor of two.

This update doesn't respect that scaling

## SGD: scaling issues

Usual solution is to “standardize” each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves “looking ahead” to compute the mean and variance, and destroys the online property of the algorithm



## SGD: scaling issues

Usual solution is to “standardize” each variable — subtract out the mean and divide by the standard deviation

$$x_j \leftarrow \frac{x_j - \text{mean}(x_j)}{\sqrt{\text{var}(x_j)}}$$

But this involves “looking ahead” to compute the mean and variance, and destroys the online property of the algorithm

Solution: The mean and variance can be updated in an online manner, in constant time, by storing auxiliary variables for each component  $j$ .

# SGD: Regularization

A “ridge” penalty  $\lambda \sum_{j=1}^d \beta_j^2$  is easily handled.

Gradient changes by an additive term  $2\lambda\beta_j$ . Update becomes

$$\begin{aligned}\beta_j &\longleftarrow \beta_j + \eta \{ (y - \pi) x_j - \lambda \beta_j \} \\ &= (1 - \eta \lambda) \beta_j + \eta (y - \pi) x_j\end{aligned}$$

$$\beta_j x_j \longleftarrow (1 - \eta \lambda) \beta_j x_j + \eta (y - \pi) x_j^2$$

Observe that this “does the right thing” whether  $\beta_j$  wants to be large positive or negative.

- *The penalty shrinks  $\beta_j$  toward zero*

# What did we learn today?

- As we penalize  $\|\beta\|^2$  from being too big, this *shrinks* the estimated coefficients toward zero.
- Ridge regression shrinks the coefficients. Good for high dimensions.
- If predictor variables are highly correlated, the model estimates may be unstable
- Stochastic gradient descent is a first order method that scales to large classification (and regression) problems
- Choosing the learning rate is a little tricky
- Difficult to parallelize, but not always necessary

# Readings

Classification is covered in Chapter 4 of our ISL book. In particular, Section 4.3 is on logistic regression, and Section 4.4 is on linear and quadratic discriminant analysis.