


S&DS 355 / 555
Introductory Machine Learning


Word Embeddings



3. litoria



4. leptodactylidae



5. rana



7. eleutherodactylus

Thursday, October 31

For Today

- Quiz ✓
- Class-based LMs (redux)
- Embeddings
- Embedding embeddings
- Next: Neural language models

Class-based bigram model

- Model takes form

$$\begin{aligned} p(w_2 | w_1) &= p(\text{class}(w_2) | \text{class}(w_1)) p(w_2 | \text{class}(w_2)) \\ &= p(c_2 | c_1) p(w_2 | c_2) \end{aligned}$$

Class-based bigram model

- Model takes form

$$\begin{aligned}p(w_2 | w_1) &= p(\text{class}(w_2) | \text{class}(w_1)) p(w_2 | \text{class}(w_2)) \\ &= p(c_2 | c_1) p(w_2 | c_2)\end{aligned}$$

- Use bottom-up agglomerative clustering to group the words.






Class-based bigram model

- Model takes form

$$\begin{aligned} p(w_2 | w_1) &= p(\text{class}(w_2) | \text{class}(w_1)) p(w_2 | \text{class}(w_2)) \\ &= p(c_2 | c_1) p(w_2 | c_2) \end{aligned}$$

- Use bottom-up agglomerative clustering to group the words.
- In each step, merge the pair of classes that gives the smallest reduction in likelihood of the data. (The MLE bigram model has the greatest likelihood.)

Sample merges



Sample clusters

Friday Monday Thursday Wednesday Tuesday Saturday Sunday weekends Sundays Saturdays
June March July April January December October November September August
people guys folks fellows CEOs chaps doubters commies unfortunates blokes
down backwards ashore sideways southward northward overboard aloft downwards adrift
water gas coal liquid acid sand carbon steam shale iron
great big vast sudden mere sheer gigantic lifelong scant colossal
man woman boy girl lawyer doctor guy farmer teacher citizen
American Indian European Japanese German African Catholic Israeli Italian Arab
pressure temperature permeability density porosity stress velocity viscosity gravity tension
mother wife father son husband brother daughter sister boss uncle
machine device controller processor CPU printer spindle subsystem compiler plotter
John George James Bob Robert Paul William Jim David Mike
anyone someone anybody somebody
feet miles pounds degrees inches barrels tons acres meters bytes
director chief professor commissioner commander treasurer founder superintendent dean cus-
todian
liberal conservative parliamentary royal progressive Tory provisional separatist federalist PQ
hadn't hath would've could've should've must've might've
asking telling wondering instructing informing kidding reminding bothering thanking deposing
that tha theat
head body hands eyes voice arm seat eye hair mouth

Group globally, compute locally

- Clusters contain syntactic and semantic elements
- Surprising, since use local statistics only
- “A word is known by the company it keeps”
- Two words are similar if they appear with similar words

Perplexity

Perplexity is a simple transformation of the likelihood:

$$\text{Perplexity}(\theta) = \left(\prod_{i=1}^N p_{\theta}(w_i | w_{1:i-1}) \right)^{-1/N}$$

Inverse of geometric mean.

If perplexity is 100, then the model predicts, on average, as if there were 100 equally likely words to follow.

Class-based bigram model

- Clusters contain a lot of information
- Each word in a single class; partition of vocabulary
- Number of parameters: $O(C^2 + C \cdot V)$ where C is number of classes.
- Reduces size of model with relatively small increase in perplexity, $244 \mapsto 271$.

Pointwise mutual information (PMI)

Average mutual information

$$I(W_1, W_2) = \sum_{w_1, w_2} p(w_1, w_2) \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

Pointwise mutual information (PMI)

$$\log \left(\frac{p_{\text{near}}(w_1, w_2)}{p(w_1)p(w_2)} \right)$$

- How likely are specific words/clusters to co-occur together within some window, compared to if they were independent?

Example clusters from PMI

we our us ourselves ours
question questions asking answer answers answering
performance performed perform performs performing
tie jacket suit
write writes writing written wrote pen
morning noon evening night nights midnight bed
attorney counsel trial court judge
problems problem solution solve analyzed solved solving
letter addressed enclosed letters correspondence
large size small larger smaller
operations operations operating operate operated
school classroom teaching grade math
street block avenue corner blocks
table tables dining chairs plate
published publication author publish writer titled
wall ceiling walls enclosure roof
sell buy selling buying sold

Shortcomings of word clusters

- Can't use vector space operations
- Doesn't give “features” of words
- These are addressed with “distributed representations” (next)

Core idea of embeddings

- Form a language model but replace classes by vectors, one for each word
- Use PMI-like scores to fit the vectors
- Can be applied whenever have cooccurrence data.

Embedding LM

Model is

$$p(w_2 | w_1) = \frac{\exp(\phi(w_2)^T \phi(w_1))}{\sum_w \exp(\phi(w)^T \phi(w_1))}.$$

As before,

$$\begin{aligned}\ell(\phi) &= \sum_{w_1, w_2} p(w_1, w_2) \log p(\phi_2 | \phi_1) p(w_2 | \phi_2) \\ &= I(\Phi_1, \Phi_2) - H(W)\end{aligned}$$

Thus, we want embedding vectors with high mutual information.

Constructing embeddings

Carry out stochastic gradient descent over the embedding vectors $\phi \in \mathbb{R}^d$ (where $d \approx 50\text{--}100$ is chosen by trial and error)

This is what Mikolov et al. (2014, 2015) did at Google. With a couple of heuristics:

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- An issue with this is that it “over generates” the data. With text
the lazy brown fox jumped **we will have** $p(\text{brown}|\text{lazy})$
and $p(\text{brown}|\text{fox})$

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- An issue with this is that it “over generates” the data. With text
the lazy brown fox jumped we will have $p(\text{brown} | \text{lazy})$
and $p(\text{brown} | \text{fox})$
- Second is computational. The bottleneck is computing the denominator in the logistic (softmax) probability.

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- An issue with this is that it “over generates” the data. With text
the lazy brown fox jumped **we** will have $p(\text{brown} | \text{lazy})$
and $p(\text{brown} | \text{fox})$
- Second is computational. The bottleneck is computing the denominator in the logistic (softmax) probability.
- Use “negative sampling”: Approximation

$$\begin{aligned} & \sum_w \exp(\phi(w)^T \phi(w_1)) \\ & \approx \exp(\phi(w_2)^T \phi(w_1)) + \sum_{\text{random } w} \exp(\phi(w)^T \phi(w_1)) \end{aligned}$$

Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

$$\phi(\text{king}) - \phi(\text{man}) \stackrel{?}{\approx} \phi(\text{queen}) - \phi(\text{woman})$$

$$\hat{w} = \arg \min_w \|\phi(\text{king}) - \phi(\text{man}) + \phi(\text{woman}) - \phi(w)\|^2$$

Does $\hat{w} = \text{queen}$?

Learned Analogies

Table 8: *Examples of the word pair relationships, using the best word vectors from Table 4 (Skip-gram model trained on 783M words with 300 dimensionality).*

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

Evaluation Analogies

Type of relationship	Word Pair 1		Word Pair 2	
Common capital city	Athens	Greece	Oslo	Norway
All capital cities	Astana	Kazakhstan	Harare	Zimbabwe
Currency	Angola	kwanza	Iran	rial
City-in-state	Chicago	Illinois	Stockton	California
Man-Woman	brother	sister	grandson	granddaughter
Adjective to adverb	apparent	apparently	rapid	rapidly
Opposite	possibly	impossibly	ethical	unethical
Comparative	great	greater	tough	tougher
Superlative	easy	easiest	lucky	luckiest
Present Participle	think	thinking	read	reading
Nationality adjective	Switzerland	Swiss	Cambodia	Cambodian
Past tense	walking	walked	swimming	swam
Plural nouns	mouse	mice	dollar	dollars
Plural verbs	work	works	speak	speaks

GloVe

Shortly after: Stanford group introduced a computational expedient (with attempt to give a “principled” motivation)

$$\mathcal{O}(\phi) = \sum_{w_1, w_2} f(c_{w_1, w_2}) \left(\phi(w_1)^T \phi(w_2) - \log c_{w_1, w_2} \right)^2$$

where $c_{w, w'}$ are cooccurrence counts.

- A type of regression estimator. Can interpret/relate this to other objectives.
- Main advantage is that SGD can be carried out much more efficiently

GloVe

$$\mathcal{O}(\phi) = \sum_{w_1, w_2} f(c_{w_1, w_2}) \left(\phi(w_1)^T \phi(w_2) - \log c_{w_1, w_2} \right)^2$$

where $c_{w, w'}$ are cooccurrence counts.

- Heuristic weighting function

$$f(x) = \left(\frac{x}{x_{\max}} \right)^\alpha$$

where $\alpha = 3/4$ set empirically.

- So $10^{-4} \mapsto 10^{-3}$. Each order of magnitude down gets “boosted” by 1/4-magnitude.

GloVe site and code

GloVe: Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher, Christopher D. Manning

Introduction

GloVe is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

Getting started (Code download)

- Download the [code](#) (licensed under the [Apache License, Version 2.0](#))
- Unpack the files: `unzip GloVe-1.2.zip`
- Compile the source: `cd GloVe-1.2 && make`
- Run the demo script, `demo.sh`
- Consult the included README for further usage details, or ask a [question](#)
- The code is also available [on GitHub](#)

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](#) v1.0 whose full text can be found at: <http://www.opensource.org/licenses/oddl/1.0/>
 - [WikiPedia 2014 - English word](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download) [glove.6B.zip](#)
 - Common Crawl (1.2B tokens, 10M vocab, uncased, 300d vectors, 175 GB download) [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 22M vocab, cased, 300d vectors, 2.03 GB download) [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 12M tokens, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download) [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data

Citing GloVe

Jeffrey Pennington, Richard Socher, and Christopher D. Manning, 2014. [GloVe: Global Vectors for Word Representation](#) ([pdf](#)) ([bib](#))

Highlights

1. Nearest neighbors

The Euclidean distance (or cosine similarity) between two word vectors provides an effective method for measuring the linguistic or semantic similarity of the corresponding words. Sometimes, the nearest neighbors according to this metric reveal rare but relevant words that lie outside an average human's vocabulary. For example, here are the closest words to the target word *frog*.

0. frog
1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



3. litoria



4. leptodactylidae



5. rana





7. eleutherodactylus

2. Linear substructures

The similarity metrics used for nearest neighbor evaluations produce a single scalar that quantifies the relatedness of two words. This simplicity can be problematic since two given words almost always exhibit more intricate relationships than can be captured by a single number. For example, men may be regarded as similar to women in that both words describe human beings; on the other hand, the two words are often considered opposites since they highlight a primary axis along which humans differ from one another.

Recommendation via Embedding



Notebook

Let's go to the Python notebook!

Embedding embeddings: t-SNE

- How can we visualize the embeddings?

Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space

Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space
- Can do PCA, but this will introduce an additional projection/approximation step

Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space
- Can do PCA, but this will introduce an additional projection/approximation step
- Many visualization techniques exist. A currently popular one is t-SNE: "Student-t Stochastic Neighborhood Embedding"

t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings

t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix $P = [P_{ij}]$

t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix $P = [P_{ij}]$
- Represent word i by $y_i \in \mathbb{R}^2$. Use a heavy-tailed distribution (Student-t with one degree of freedom)

t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix $P = [P_{ij}]$
- Represent word i by $y_i \in \mathbb{R}^2$. Use a heavy-tailed distribution (Student-t with one degree of freedom)
- Select y_i using stochastic gradient descent

t-SNE: Detailed algorithm

For each word w_i compute a language model

$$P_{j|i} \propto \exp \left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

That is:

$$P_{j|i} = \frac{\exp \left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)}{\sum_k \exp \left(-\frac{\|\phi(w_i) - \phi(w_k)\|^2}{2h_i^2} \right)}$$

t-SNE: Detailed algorithm

For each word w_i compute a language model

$$P_{j|i} \propto \exp \left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

That is:

$$P_{j|i} = \frac{\exp \left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)}{\sum_k \exp \left(-\frac{\|\phi(w_i) - \phi(w_k)\|^2}{2h_i^2} \right)}$$

Choose the bandwidth h_i so that the perplexity is, say, 10. This puts the probabilities all on the same scale.

t-SNE: Detailed algorithm

For each word w_i compute a language model

$$P_{j|i} \propto \exp\left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2}\right)$$

t-SNE: Detailed algorithm

For each word w_i compute a language model

$$P_{j|i} \propto \exp \left(-\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

Now form

$$P_{ij} = \frac{1}{2} (P_{j|i} + P_{i|j})$$

as a simple way of symmetrizing.

t-SNE: Detailed algorithm

Now form Student-t distribution depending on the visualization vectors $y_i \in \mathbb{R}^2$:

$$Q_{ij} \propto \left(1 + \|y_i - y_j\|^2\right)^{-1}$$

t-SNE: Detailed algorithm

Now form Student-t distribution depending on the visualization vectors $y_i \in \mathbb{R}^2$:

$$Q_{ij} \propto \left(1 + \|y_i - y_j\|^2\right)^{-1}$$

That is:

$$Q_{ij} = \frac{\left(1 + \|y_i - y_j\|^2\right)^{-1}}{\sum_{k \neq \ell} \left(1 + \|y_k - y_\ell\|^2\right)^{-1}}$$

This has fatter tails than a Gaussian

t-SNE: Detailed algorithm

Finally, run stochastic gradient descent (SGD) over the vectors y_i to optimize:

$$\begin{aligned}\hat{y} &= \arg \min \sum_{ij} P_{ij} \log P_{ij} / Q_{ij} \\ &= \arg \max \sum_{ij} P_{ij} \log Q_{ij}\end{aligned}$$

Interpretation: if $\phi(w_i)$ is very close to $\phi(w_j)$ then y_i will be close to y_j .
(long distances may be stretched further...)


t-SNE: More info and examples

<https://lvdmaaten.github.io/tsne/>

<http://cs.stanford.edu/people/karpathy/tsnejs/>

Note: This is just a visualization technique, to give intuition for the high dimensional embedding

t-SNE: Examples



<http://www.cs.cornell.edu/~ginsparg/arxiv/gmaps2.html>

Summary: Word embeddings

- Word embeddings are vector representations of words, learned from cooccurrence statistics
- The models can be viewed in terms of logistic regression and class-based bigram models
- Surprising semantic relations are encoded in linear relations
- Various heuristics have been introduced to get scalability
- Embeddings improve with more data
- t-SNE is an algorithm for visualizing embeddings